

# **Brain Inspired Artificial Intelligence**

## **2: Markov Decision Process**


Eiji Uchibe

Dept. of Brain Robot Interface

ATR Computational Neuroscience Labs.


# Reminder: Bandit Problem

- Bandit problems: one state, many actions
- The goal is to find an optimal policy  $\pi(a)$  that maximize the expected reward
- The value based approach estimates the value function  $q_{\pi}(a)$
- The policy based approach directly search the optimal policy
- Now, we will study Markov Decision Processes (MDPs) for sequential decision making
  - MDPs formally describe an environment for reinforcement learning






# Extension to Sequential Decision Making Problems

- A learning agent is not in front of gambling machines
- To get non-zero rewards, the agent has to move to one of machines
- A current position is considered as a state
  - The state is a sufficient statistics to describe the dynamics of the environment
  - Bandits are MDPs with one state
- The action of the agent influences the environment, causing a state transition

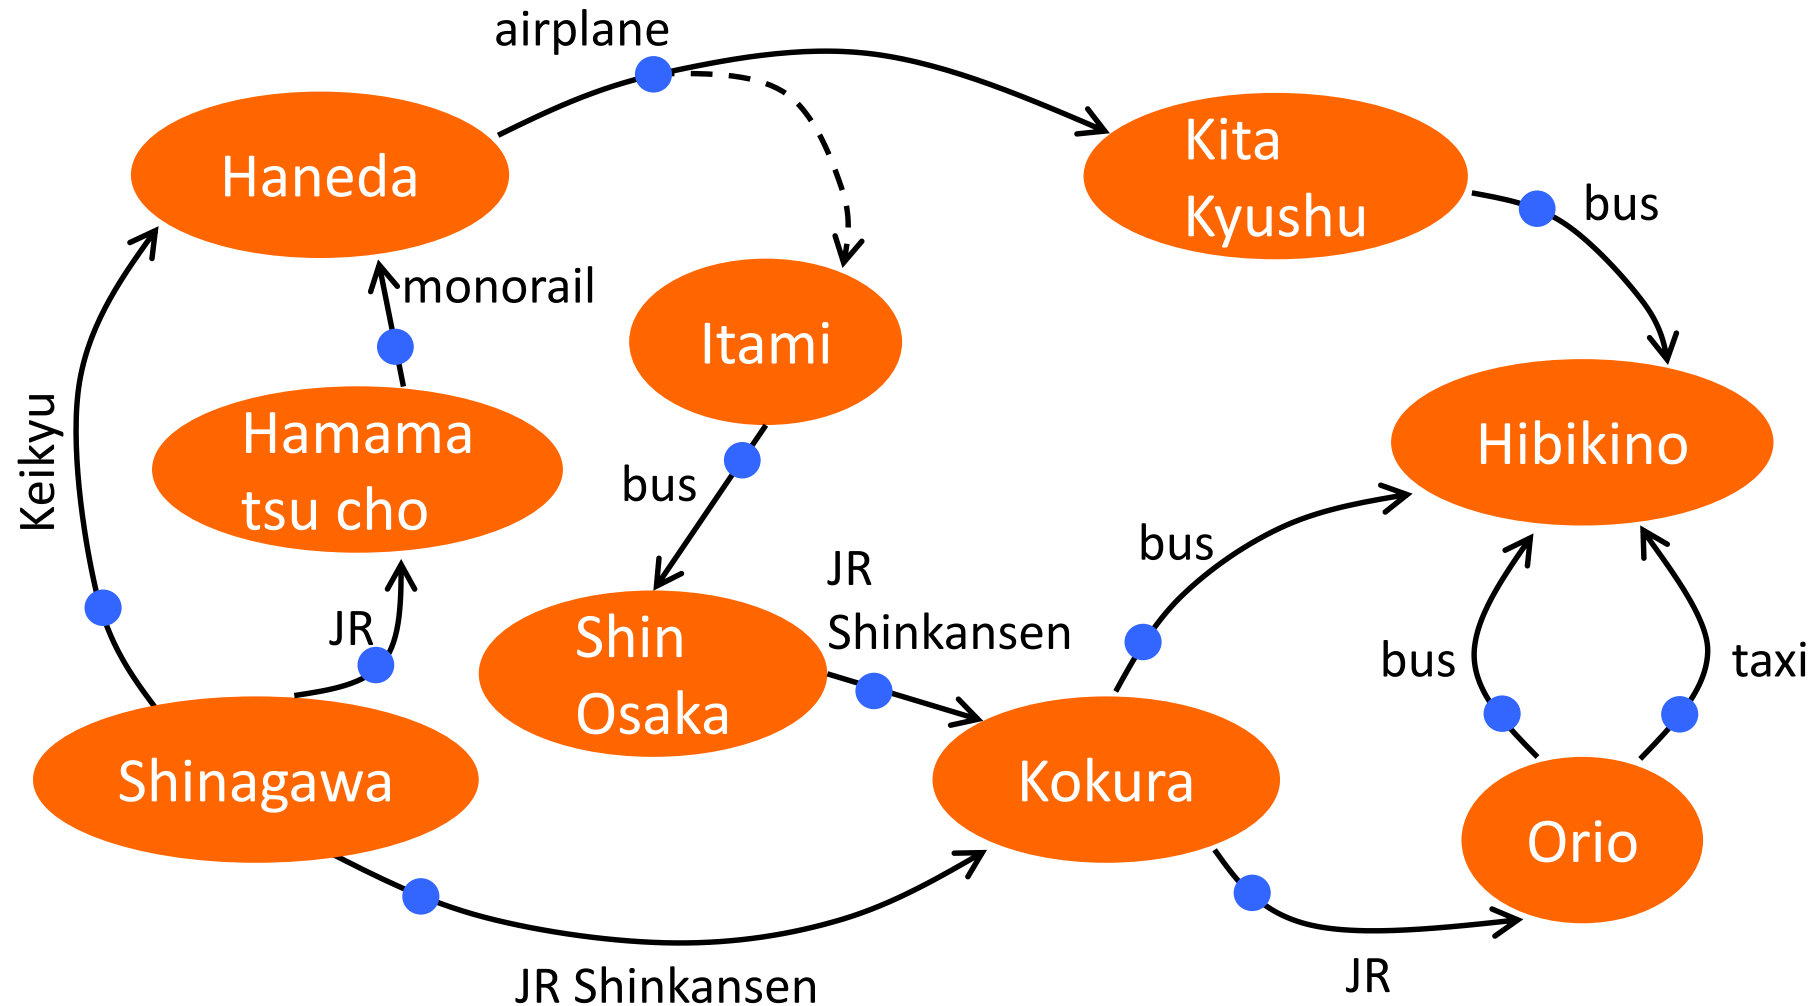


# Extension to the Sequential Decision Making Problems

- $s$ : state, 
- $a$ : action, 
- State set:  $\mathcal{S} = \{s_1, s_2, \dots, s_{|\mathcal{S}|}\}$
- Action set:  $\mathcal{A} = \{a_1, a_2, \dots, a_{|\mathcal{A}|}\}$
- $p_0(s)$ : initial state distribution
- $p(s', r \mid s, a)$ : stochastic environmental dynamics
- $\pi(a \mid s)$ : policy, probability to select an action  $a$  for state  $s$

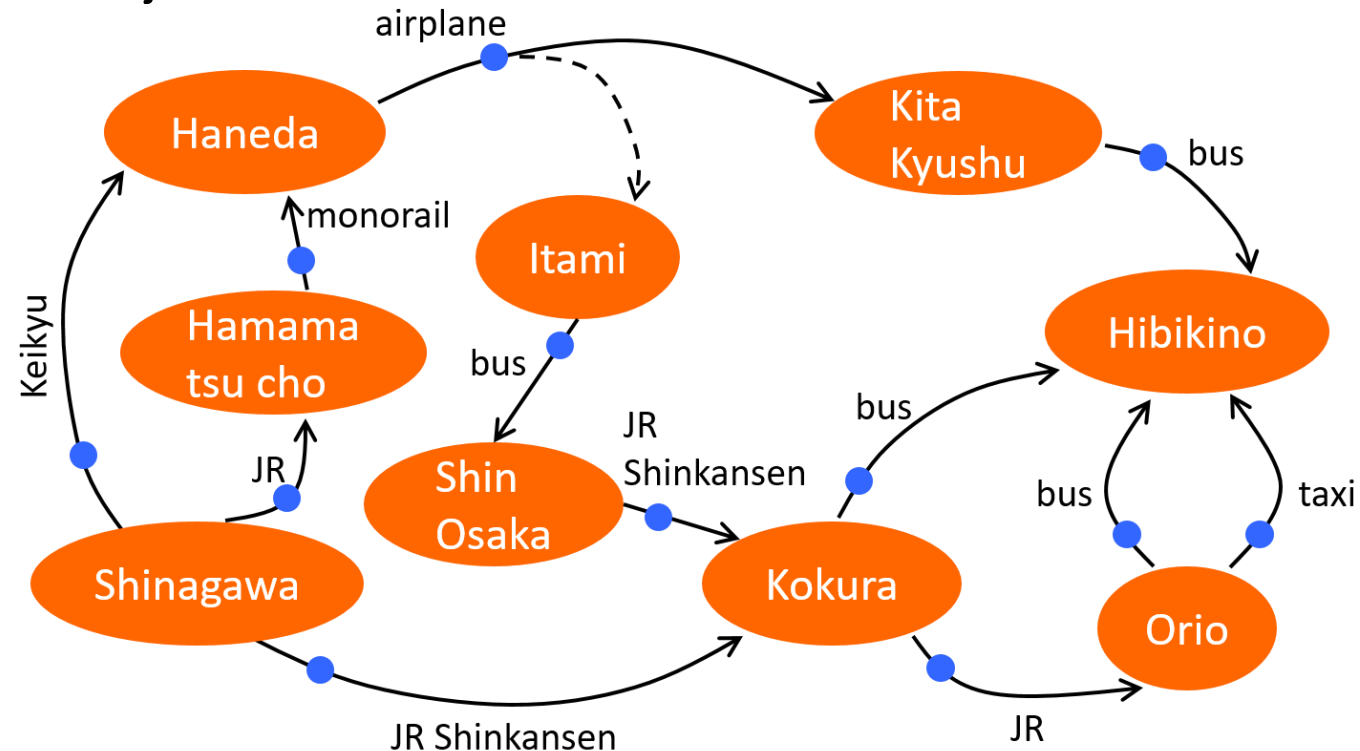


# Example: From Shinagawa to Hibikino Campus



# States and Actions in the Example

- $\mathcal{S} = \{\text{Shinagawa, Hamamatsu cho, Haneda, Itami, Shin Osaka, Kitakyushu, Kokura, Orio, Hibikino}\}$
- Action set is state dependent in this case
  - $\mathcal{A}(\text{Shinagawa}) = \{\text{Keikyu, JR Shinkansen}\}$
  - $\mathcal{A}(\text{Kokura}) = \{\text{bus, JR Shinkansen}\}$
  - $\mathcal{A}(\text{Shinagawa}) = \{\text{Airplane, JR Shinkansen}\}$
  - $\vdots$



# Markovian Assumption



- Consider a sequence of states, action, and rewards from  $t = 0$  to  $t$

$$S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{t-1}, A_{t-1}, R_t, S_t, A_t$$

- How can we model the next state  $S_{t+1}$  and the reward  $R_{t+1}$ ?  
The most general way is to introduce a conditional probability distribution


$$\Pr(S_{t+1} = s', R_{t+1} = r \mid S_0, A_0, R_1, \dots, S_{t-1}, A_{t-1}, R_t, S_t, A_t)$$

- It is too complicated! We usually **assume** that the environment satisfies the following Markov property:

$$p(s', r \mid s, a) = \Pr(S_{t+1} = s', R_{t+1} = r \mid S_t = s, A_t = a)$$

# State Representation is Important

- Consider a stochastic state transition
  - $\Pr(\text{Kita Kyushu} | \text{Haneda, airplane}) = 0.9$
  - $\Pr(\text{Itami} | \text{Haneda, airplane}) = 0.1$
- Markovian property
  - $\Pr(\text{Kita Kyushu} | \text{Haneda, airplane})$   
=  $\Pr(\text{Kita Kyushu} | \text{Haneda, airplane, Shinagawa, Keikyu})$   
=  $\Pr(\text{Kita Kyushu} | \text{Haneda, airplane, Hamamatsu cho, monorail, Shinagawa, JR})$





# State Representation is Important

- Mario's position  $\mathbf{p}_t$  at time  $t$  is not uniquely determined if we use the game screen as state

$$\Pr(\mathbf{p}_{t+1} \mid \mathbf{p}_t, a_t) \neq \Pr(\mathbf{p}_{t+1} \mid \mathbf{p}_t, a_t, \mathbf{p}_{t-1})$$

- For example, one defines the state by

$$s_t = (\mathbf{p}_t, \mathbf{p}_{t-1})$$


- $s_t$  contains the information of the velocity of Mario

- If a state is given by the position and velocity of the cart and the angle and angular velocity of the pole, it satisfies Markov assumption

$$s_t = (x_t, \dot{x}_t, \theta, \dot{\theta})$$




Mario benchmark



cart-pole system


# Discretizing State Space

- When a state is given by a continuous vector, it is discretized manually
  - Example of the soccer robot
  - Task is to shoot a ball into a goal
  - 5 state variables:
    - x-position and size of the ball and
    - x-position, size and orientation of the goal
- The number of states grows exponentially as the number of features increases




**Ball**

position



left center right


size



small medium large


**Goal**

position




left center right


size



small




medium




large


orientation



left-oriented



front



right-oriented

$$3^2 \times 3^3 = 243$$

the number of  
ball states

the number of  
goal states

the total number  
of states

# Components

- State transition probability:  $p(s' | s, a) = \sum_r p(s', r | s, a)$

- Expected reward for state-action pairs:

$$r(s, a) = \sum_r r p(r | s, a) \quad \text{where} \quad p(r | s, a) = \sum_{s'} p(s', r | s, a)$$

- Expected reward for state-action-next-state triples


$$r(s, a, s') = \sum_r r p(r | s, a, s') \quad \text{where} \quad p(r | s, a, s') = \frac{p(s', r | s, a)}{p(s' | s, a)}$$

# Return

- Consider a sequence of states, action, and rewards
- The return  $G_t$  is defined as the sum of discounted rewards from time-step  $t$

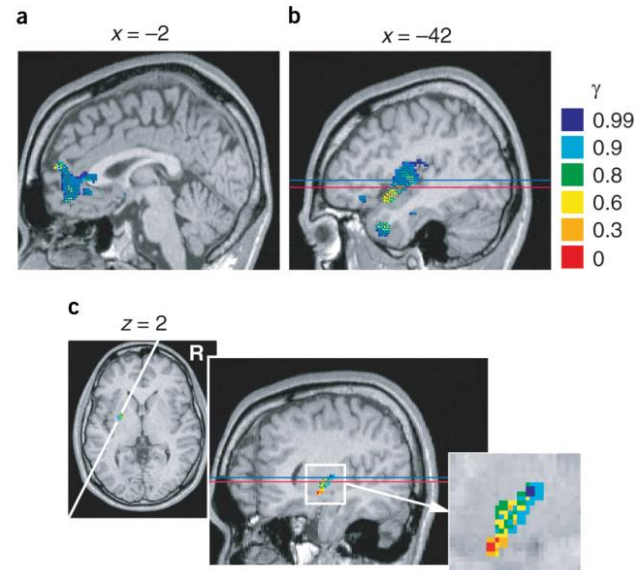
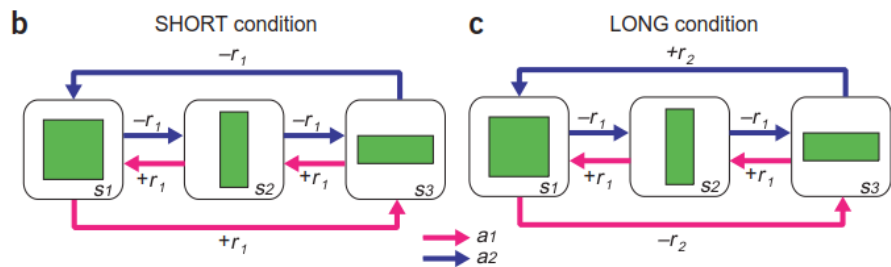

$$\begin{aligned} G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots \\ &= \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \end{aligned}$$

- $G_t$  is a random variable
- $\gamma \in [0, 1)$  is the discount rate
  - $\gamma$  close to 0: myopic evaluation
  - $\gamma$  close to 1: far-sighted evaluation




# Why Discount?

- If the reward is bounded, the return is also bounded:  
 $|R| \leq R_{\max} \quad \longrightarrow \quad |G| \leq \frac{R_{\max}}{1 - \gamma}$
- Prediction of immediate and future rewards differentially recruits cortico-basal ganglia loops



The robot does not move towards the battery



The robot tries to catch the battery

# State Value Function


- A state value function evaluates the policy in each state
- For a given stationary policy  $\pi$ , consider a state-action-reward sequence

$$S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}, R_{t+2}, S_{t+2} \dots$$

generated by

- $A_t \sim \pi(a \mid S_t)$
- $S_{t+1}, R_{t+1} \sim p(s', r \mid S_t, A_t)$
- The state value function is the expected return given by

$$v_{\pi}(s) \triangleq \mathbb{E}_{\pi}[G_t \mid S_t = s, A_{t:\infty} \sim \pi]$$



# State Value Function

- The value of a state  $s$  under a policy  $\pi$  is defined by

$$v_{\pi}(s) \triangleq \mathbb{E}_{\pi} \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s \right]$$

- $\mathbb{E}_{\pi}\{ \quad \}$  denotes the expectation which is taken over the probability distribution

$$\begin{aligned} &P(A_t, R_{t+1}, S_{t+1}, A_{t+1}, R_{t+2}, S_{t+2}, \dots \mid S_t = s) \\ &= \prod_{k=0}^{\infty} \pi(A_{t+k} \mid S_{t+k}) P_T(S_{t+k+1}, R_{t+k+1} \mid S_{t+k}, A_{t+k}) \end{aligned}$$

- Therefore,


$$v_{\pi}(s) = \sum_{A_t, R_{t+1}, S_{t+1}, \dots} P(A_t, R_{t+1}, S_{t+1}, \dots \mid S_t = s) \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \right]$$

# Example of State Value Function

- Task: to get the battery pack while avoiding collisions with an obstacle
  - Positive reward for successful battery catching
  - Negative reward for collisions with obstacles
  - Small negative reward for every time step



The goal state has the highest value




state value function



# State-Action Value Function

- A state-action value function used in the bandit problem is extended to a state-action value function
- It evaluates the all actions in each state
- It is the expected return starting from  $s$ , taking the action  $a$ , and **thereafter following policy  $\pi$** 
  - $S_{t+1}, R_{t+1} \sim p(s', r | S_t, A_t)$
  - $A_{t+1} \sim \pi(a | S_{t+1})$
- The state-action value function is the expected return given by

$$q_{\pi}(s, a) \triangleq \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$$



# State-Action Value Function

- Similarly, the value of taking action  $a$  in a state  $s$  under a policy  $\pi$  is defined by

$$q_{\pi}(s, a) \triangleq \mathbb{E}_{\pi} \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a \right]$$

- $\mathbb{E}_{\pi} \{ \quad \}$  denotes the expectation which is taken over the probability distribution

$$\begin{aligned} &P(R_{t+1}, S_{t+1}, A_{t+1}, S_{t+2}, \dots \mid S_t = s, A_t = a) \\ &= \prod_{k=0}^{\infty} p(S_{t+k+1}, R_{t+k+1} \mid S_{t+k}, A_{t+k}) \pi(A_{t+k+1} \mid S_{t+k+1}) \end{aligned}$$

- Note the difference between the distributions
- Therefore,

$$q_{\pi}(s, a) = \sum_{R_{t+1}, S_{t+1}, A_{t+1}, \dots} P(R_{t+1}, S_{t+1}, \dots \mid S_t = s, A_t = a) \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \right]$$

# Relations between $v_\pi$ and $q_\pi$

- Relation between  $v_\pi(s)$  and  $q_\pi(s, a)$

$$\begin{aligned} v_\pi(s) &= \sum_{a \in \mathcal{A}} \pi(a | s) q_\pi(s, a) \\ &= \mathbb{E}_\pi[q_\pi(s, a)] \end{aligned}$$



- Advantage function

$$A_\pi(s, a) \triangleq q_\pi(s, a) - v_\pi(s)$$

– From the definition,  $\mathbb{E}_\pi[A_\pi(s, a)] = 0$

- Relation between  $q_\pi(s, a)$  and  $v_\pi(s')$

$$q_\pi(s, a) = \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]$$



# Bellman (Expectation) Equation


- Recursive relationship of values between the current and the next state
- **Linear** with respect to  $v_\pi$
- For any policy  $\pi$  and any state  $s$ ,

$$v_\pi(s) = \mathbb{E}_\pi[G_t \mid S_t = s]$$

$$= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s]$$

$$= \sum_a \pi(a \mid s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma \mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s']]$$

$$= \sum_a \pi(a \mid s) \sum_{s', r} P(s', r \mid s, a) [r + \gamma v_\pi(s')]$$




# Bellman (Expectation) Equation

- Similarly, we obtain the recursive relationships for  $q_\pi$
- **Linear** with respect to  $q_\pi$
- For any policy  $\pi$  and any state-action pair  $(s, a)$ ,

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a]$$

$$= \sum_{s', r} p(s', r \mid s, a) [r + \gamma \mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s', A_{t+1} = a']]$$

$$= \sum_{s', r} P(s', r \mid s, a) \left[ r + \gamma \sum_{a'} \pi(a' \mid s') q_\pi(s', a') \right]$$



# Optimal Value Functions

- Reminder: The goal of RL is to find an optimal policy  $\pi_*$
- $v_*$  is an optimal state value function for  $\pi^*$  defined by

$$v_*(s) \triangleq \max_{\pi} v_{\pi}(s) \quad \forall s \in \mathcal{S}$$

- Similarly, an optimal state-action value function is defined by

$$q_*(s, a) \triangleq \max_{\pi} q_{\pi}(s, a) \quad \forall s \in \mathcal{S}, \forall a \in \mathcal{A}$$

# Theorem

- A value function defines a partial ordering over policies

$$\pi \geq \pi' \quad \text{if} \quad v_{\pi}(s) \geq v_{\pi'}(s), \quad \text{for all } s$$

- For any Markov Decision Process

1. There exists an optimal policy  $\pi_*$  that is better than or equal to all other policies,  $\pi_* \geq \pi$  for all  $\pi$
2. All optimal policies achieve the optimal value function,

$$v_{\pi_*}(s) = v_*(s)$$

3. All optimal policies achieve the optimal state-action value function

$$\underbrace{q_{\pi_*}(s, a)}_{\text{state-action value function of the optimal policy}} = \underbrace{q_*(s, a)}_{\text{optimal state-action value function}}$$

state-action value

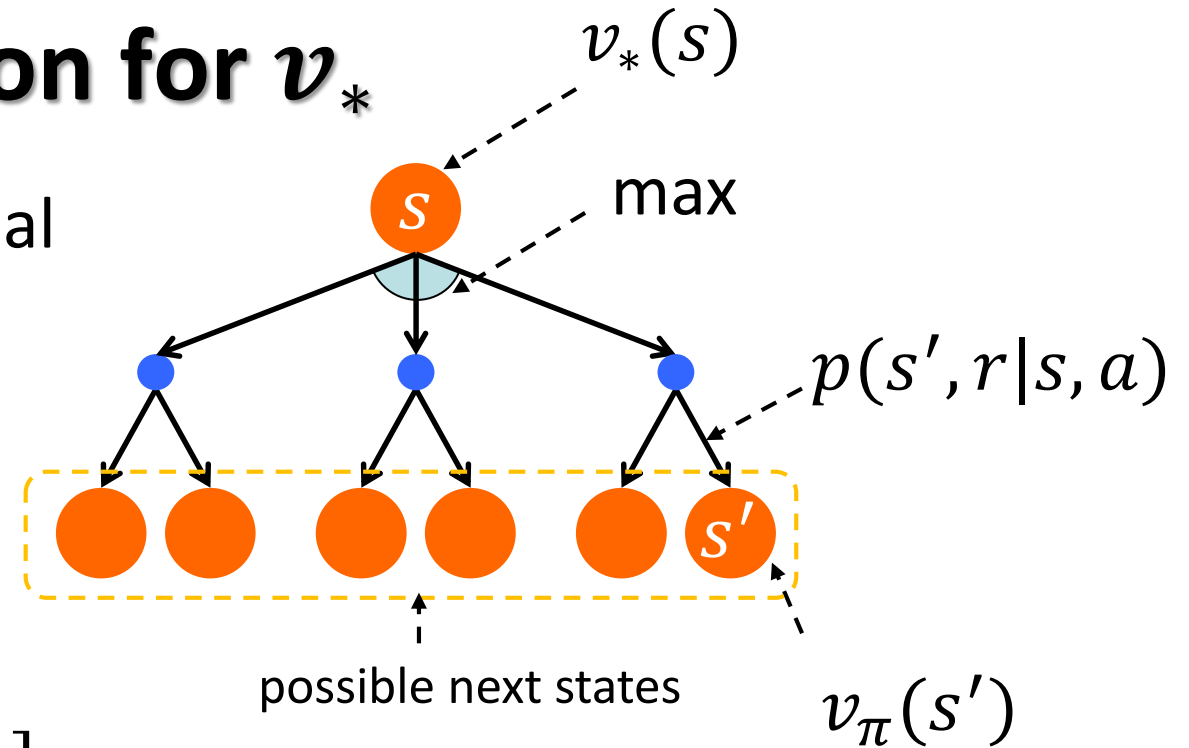
function of the optimal policy

optimal state-action  
value function

# Bellman Optimality Equation for $v_*$

- Recursive relationships for the optimal state value function
- **Nonlinear** with respect to  $v_*$
- For any state  $s$ ,

$$\begin{aligned} v_*(s) &= \max_a q_*(s, a) \\ &= \max_a \mathbb{E}_{\pi_*} [G_t \mid S_t = s, A_t = a] \\ &= \max_a \mathbb{E}_{\pi_*} [R_{t+1} + \gamma G_t \mid S_t = s, A_t = a] \\ &= \max_a \mathbb{E} [R_{t+1} + \gamma v_*(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s')] \end{aligned}$$





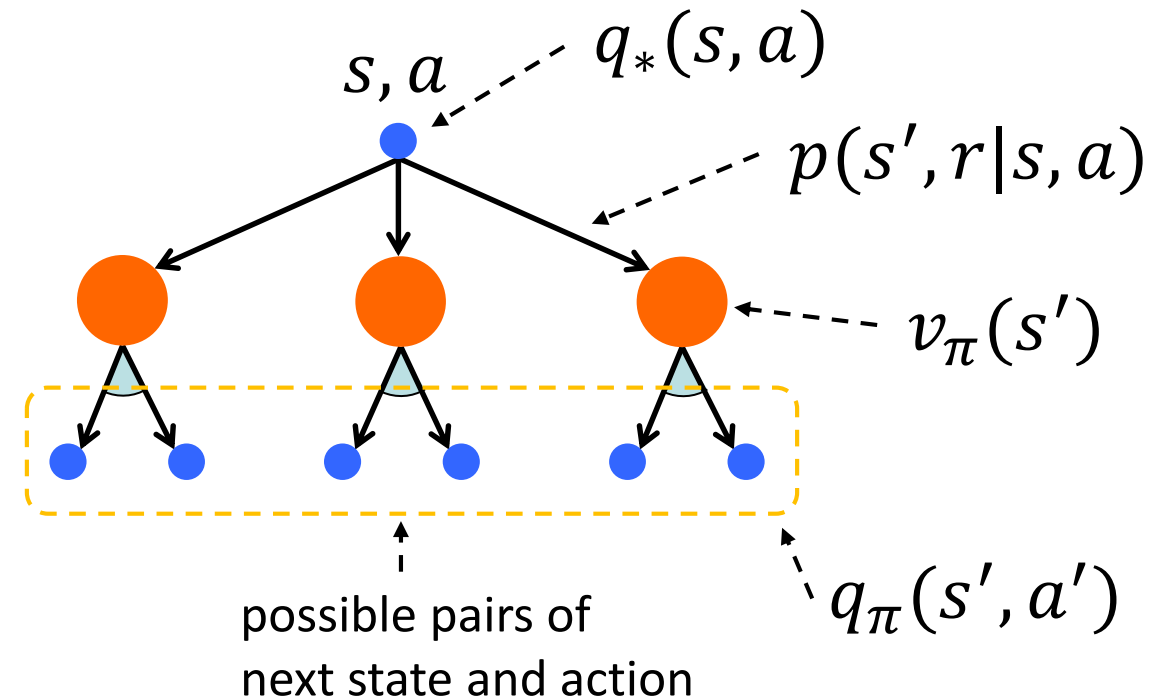
# Bellman Optimality Equation for $q_*$

- Similarly, the Bellman optimality equation is given by

$$q_*(s, a) = \mathbb{E} \left[ R_{t+1} + \gamma \max_{a'} q_*(S_{t+1}, a') \mid S_t = s, A_t = a \right]$$

$$= \sum_{s', r} p(s', r \mid s, a) \left[ r + \gamma \max_{a'} q_*(s', a') \right]$$

- This is also a nonlinear equation with respect to  $q_\pi$  due to the max operator



# Optimal Value Functions

- There exists the following relationships:

$$v_*(s) = \mathbb{E} \left[ R_{t+1} + \gamma \max_a q_*(S_{t+1}, a) \mid S_t = s \right]$$

$$q_*(s, a) = \mathbb{E}[R_{t+1} + \gamma v_*(S_{t+1}) \mid S_t = s, A_t = a]$$

# Summary

- Bellman expectation equation

$$v_{\pi}(s) = \sum_a \pi(a | s) \sum_{s', r} P(s', r | s, a) [r + \gamma v_{\pi}(s')]$$

$$q_{\pi}(s, a) = \sum_{s', r} P(s', r | s, a) \left[ r + \gamma \sum_{a'} \pi(a' | s') q_{\pi}(s', a') \right]$$

- Bellman optimality equation

$$v_*(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]$$

$$q_*(s, a) = \sum_{s', r} p(s', r | s, a) \left[ r + \gamma \max_{a'} q_*(s', a') \right]$$

# Break!



# References

- Asada, M., Noda, S., Tawaratsumida, S., and Hosoda, K. (1996). Purposive behavior acquisition for a real robot by vision-based reinforcement learning. Machine Learning, 23, 279-303, 1996.
- Playing Mario with Deep Reinforcement Learning. <https://github.com/aleju/mario-ai>
- Tanaka, S.C., Doya, K., Okada, G., Ueda, K., Okamoto, Y., and Yamawaki, S. Prediction of immediate and future rewards differentially recruits cortico-basal ganglia loops. Nature Neuroscience, 7(8): 887-893, 2004.