

Brain-Inspired Artificial Intelligence

2: Markov Decision Process


Eiji Uchibe

Dept. of Brain Robot Interface

ATR Computational Neuroscience Labs.

Reminder: Bandit Problem

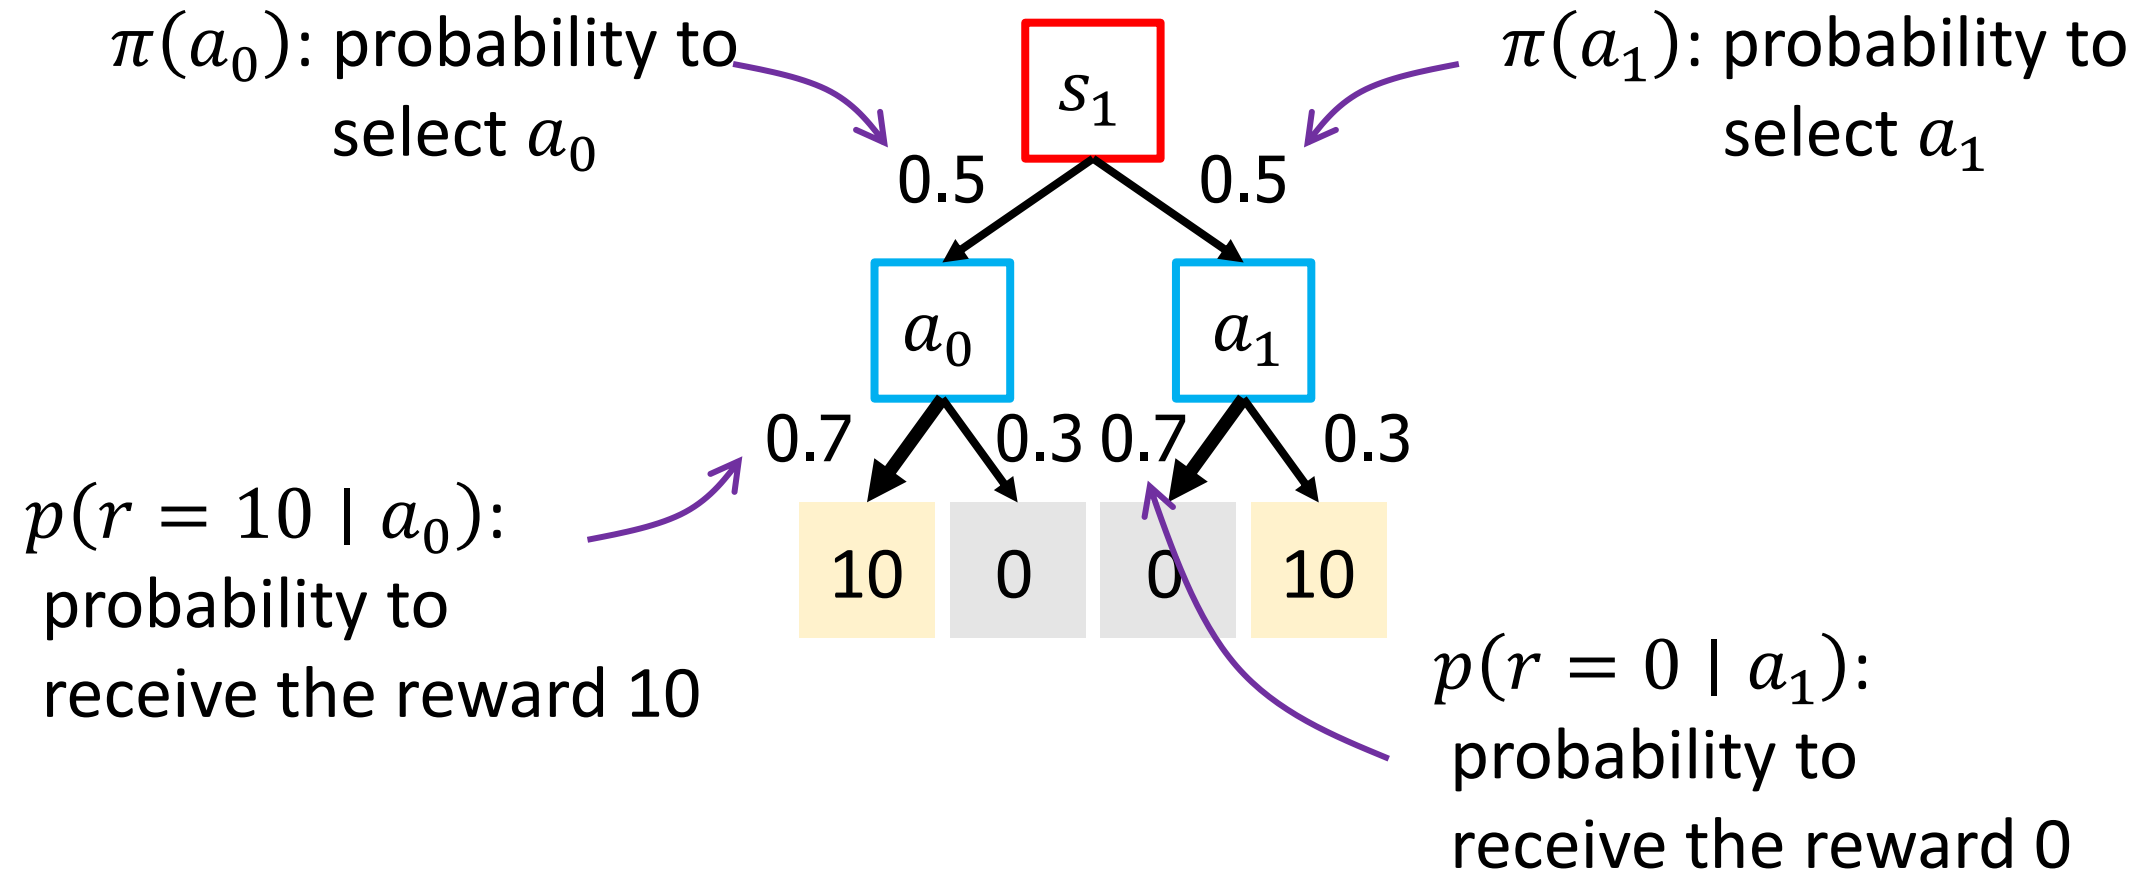
- One state, many actions
 - The goal is to find an optimal policy $\pi(a)$ that maximize the expected reward
 - The value based approach estimates the value function $Q^\pi(a)$
 - The policy based approach directly search the optimal policy
- Now, we will study Markov Decision Process (MDP) for sequential decision making
 - MDP formally describes an environment for reinforcement learning



ONE-STAGE MARKOV DECISION PROCESS

One-stage MDP

- 1-state 2-action task




Compute the action-values

- Expected value of reward

$$\begin{aligned} - Q(a_0) &= \sum_r rp(r | a_0) \\ &= 10 \times 0.7 + 0 \times 0.3 = 7 \end{aligned}$$

$$- Q(a_1) = 0 \times 0.7 + 10 \times 0.3 = 3$$



Compute the action-values


- Expected value of reward

$$\begin{aligned} - Q(a_0) &= \sum_r rp(r | a_0) \\ &= 10 \times 0.7 + 0 \times 0.3 = 7 \end{aligned}$$

$$- Q(a_1) = 0 \times 0.7 + 10 \times 0.3 = 3$$

- More formally,

$$- Q(a) = \mathbb{E}_{r \sim p(r|a)}[r]$$




Compute the state-value

- The value of state s_1 depends on the policy.
Suppose the policy is given by

$$\pi(a_0) = \pi(a_1) = 0.5.$$


- Possible transitions
 - $a_0 \rightarrow 10$ with probability 0.5×0.7
 - $a_0 \rightarrow 0$ with probability 0.5×0.3
 - $a_1 \rightarrow 0$ with probability 0.5×0.7
 - $a_1 \rightarrow 10$ with probability 0.5×0.3



$$\begin{aligned} V(s_1) &= 10 \times 0.5 \times 0.7 \\ &\quad + 10 \times 0.5 \times 0.3 \\ &= 5 \end{aligned}$$

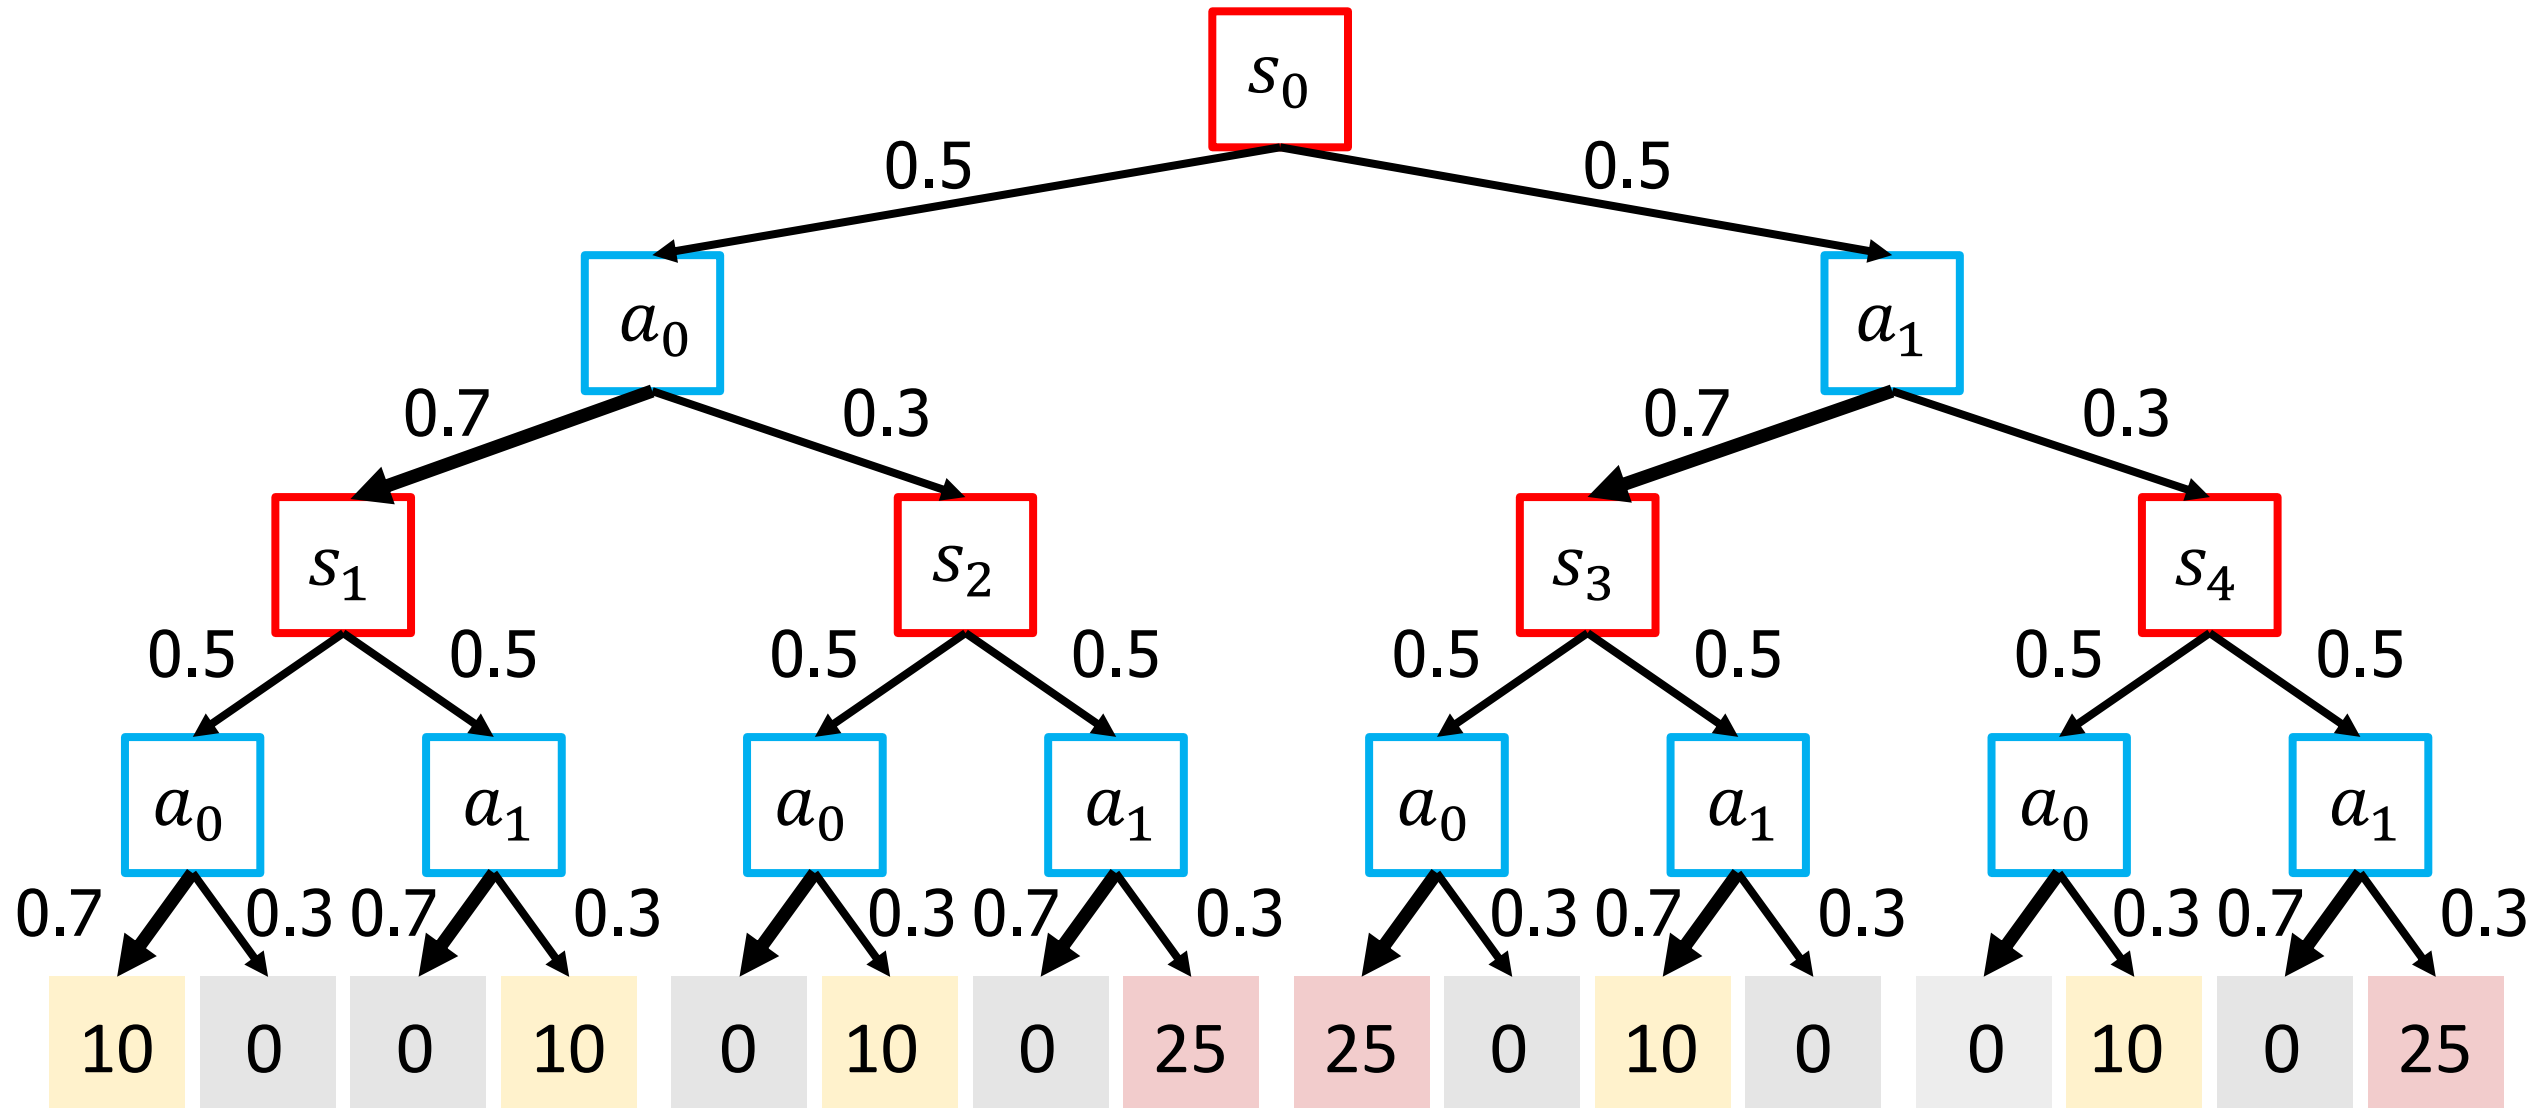
Compute the state-value from the action-values

- The action values $Q(a_1)$ and $Q(a_2)$ can be used to compute the value of state
- $$V(s_1) = \mathbb{E}_{a \sim \pi(a)}[Q(a)] = \sum_a Q(a)\pi(a)$$
$$= 7 \times 0.5 + 3 \times 0.5 = 5$$
- Hereafter, to **clarify the dependency of π** , the state-value function is denoted by $V^\pi(s)$

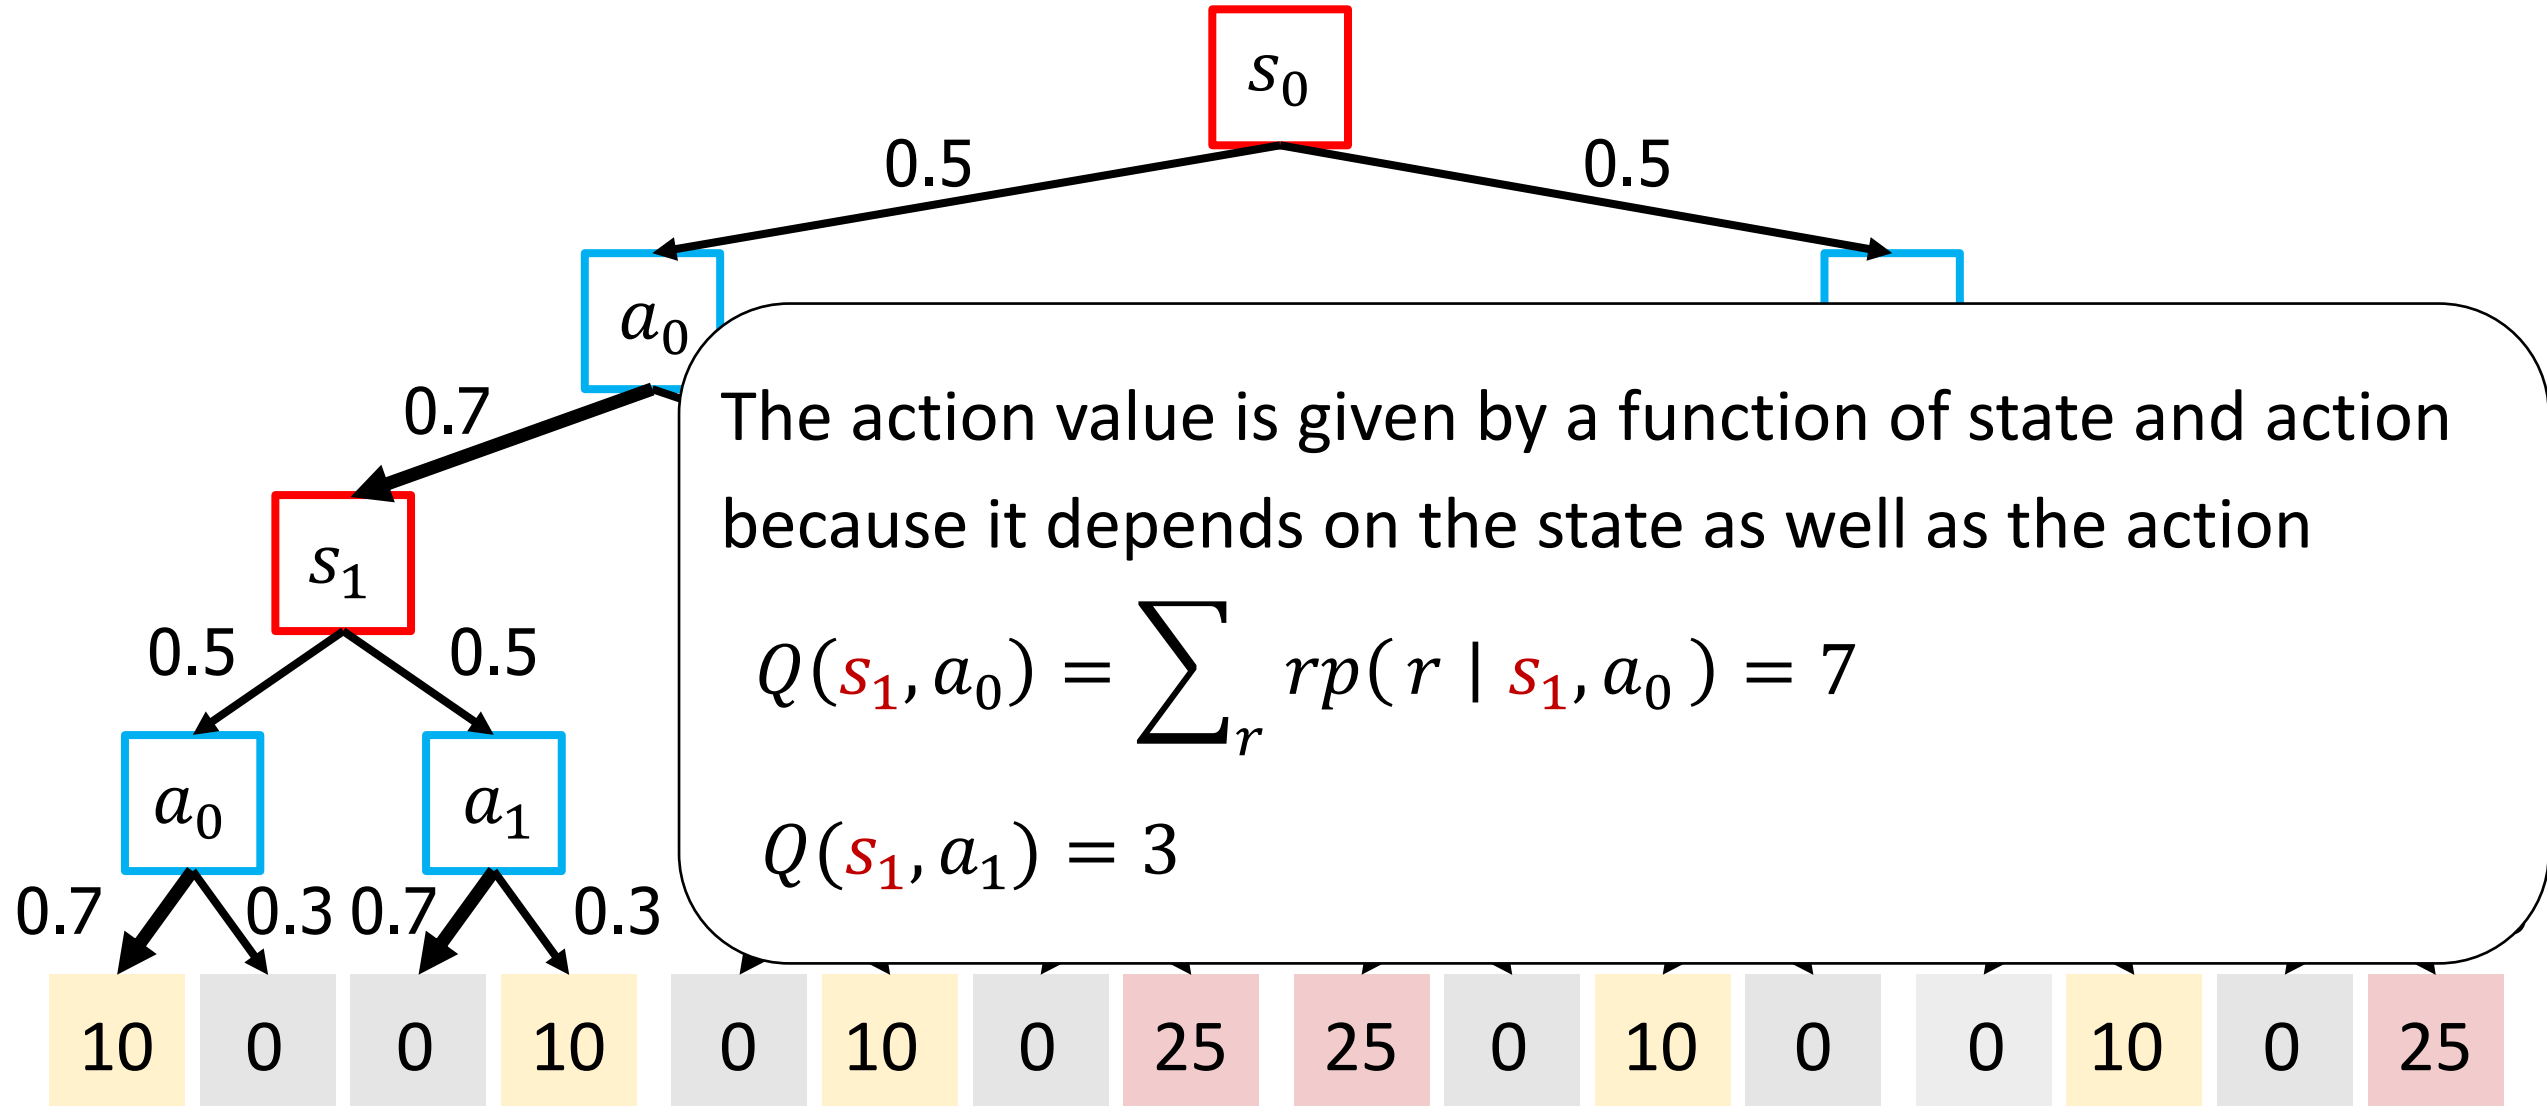


TWO-STAGE MARKOV DECISION PROCESS

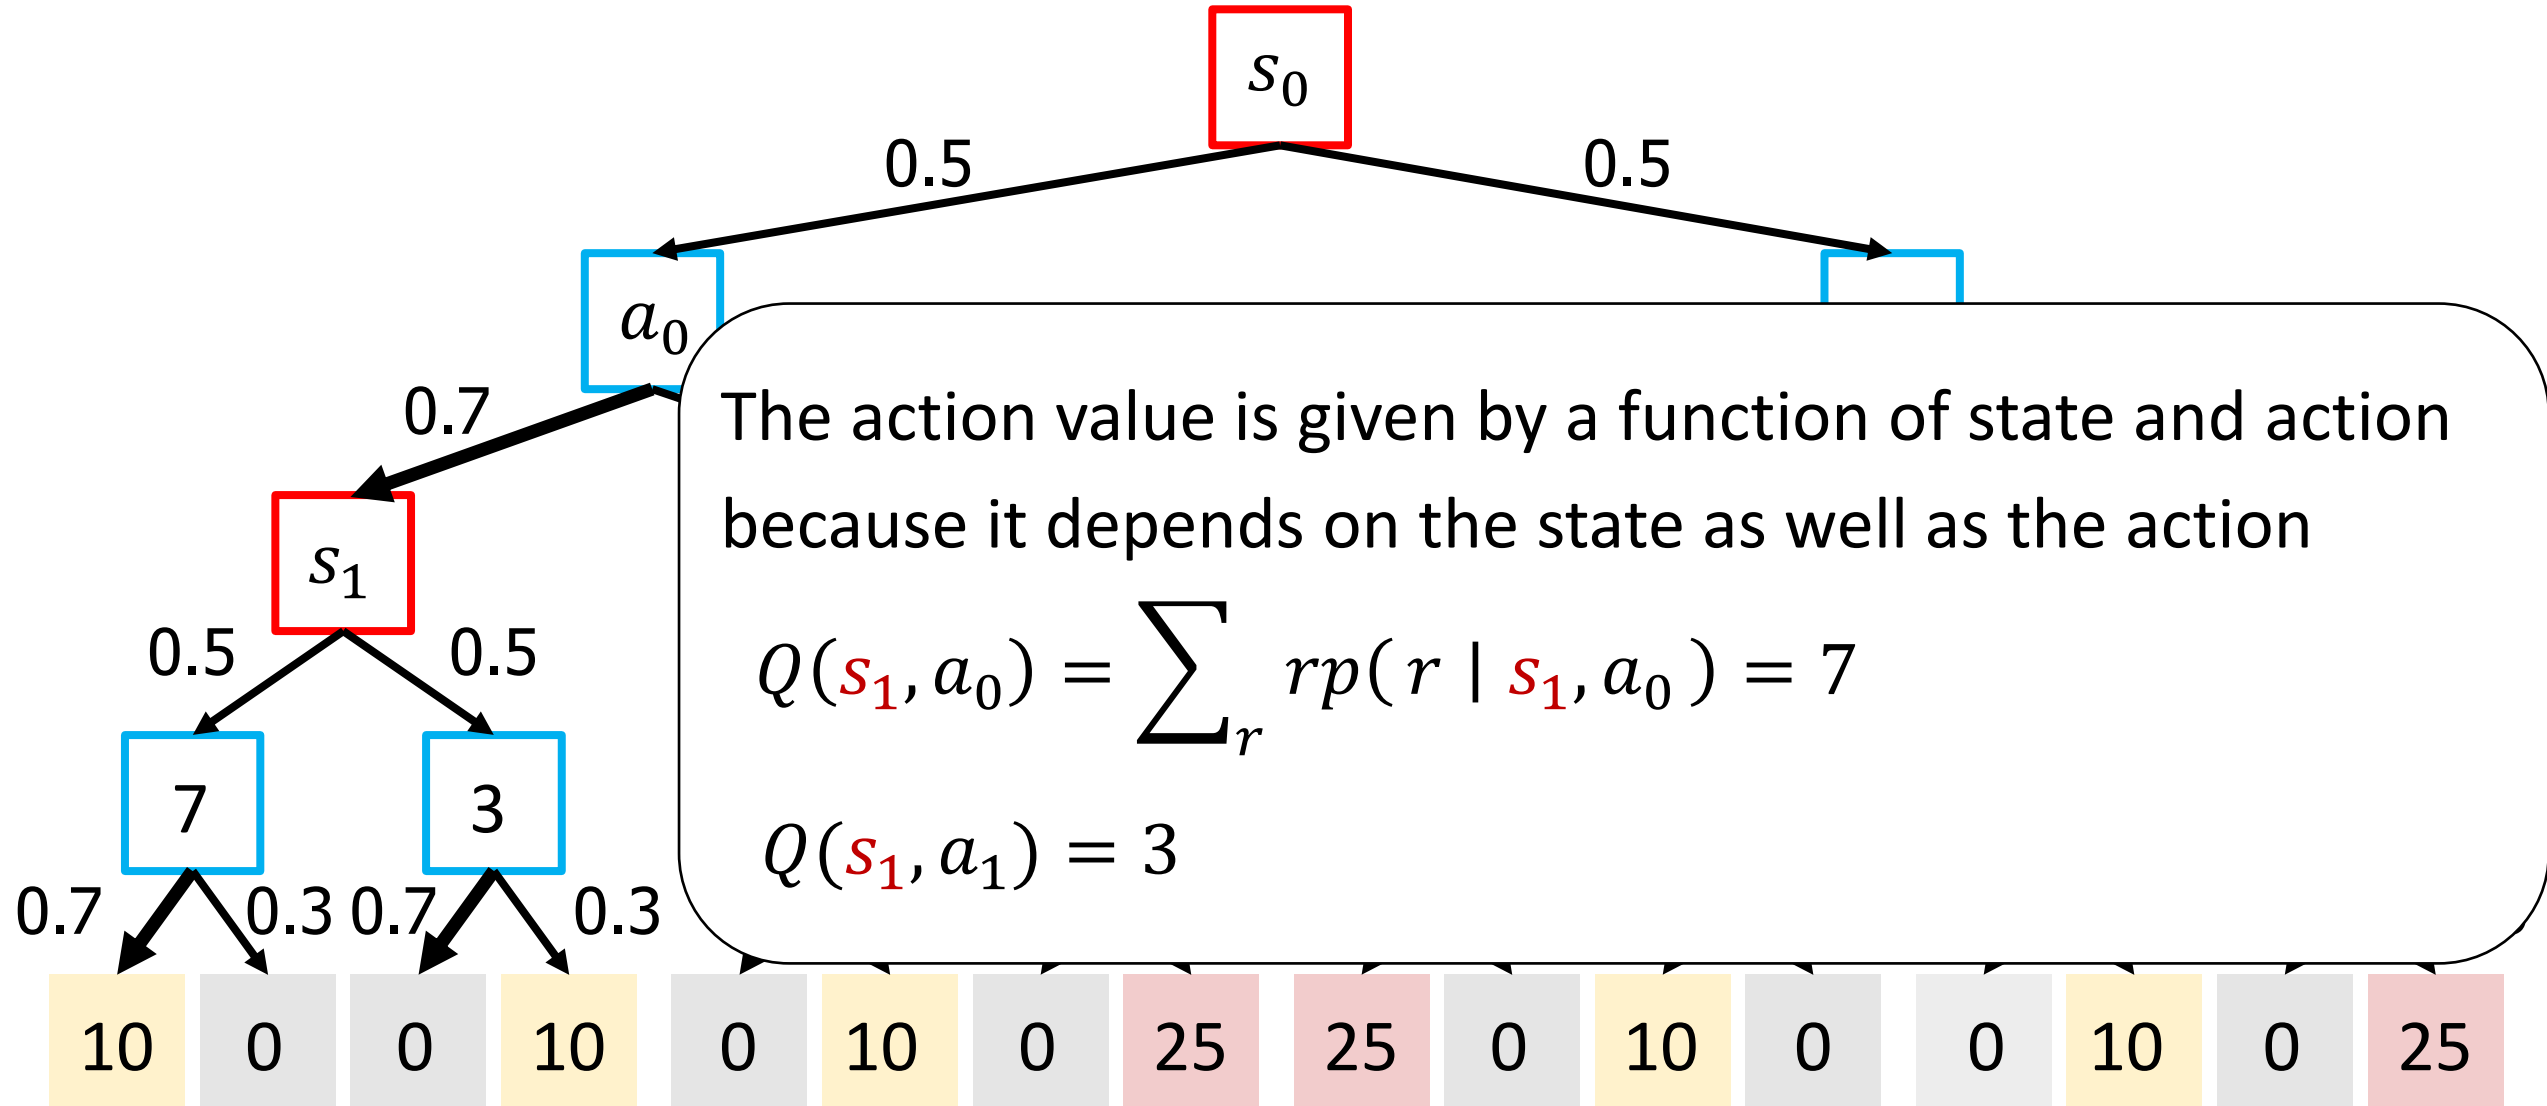
Extension to a two-stage task



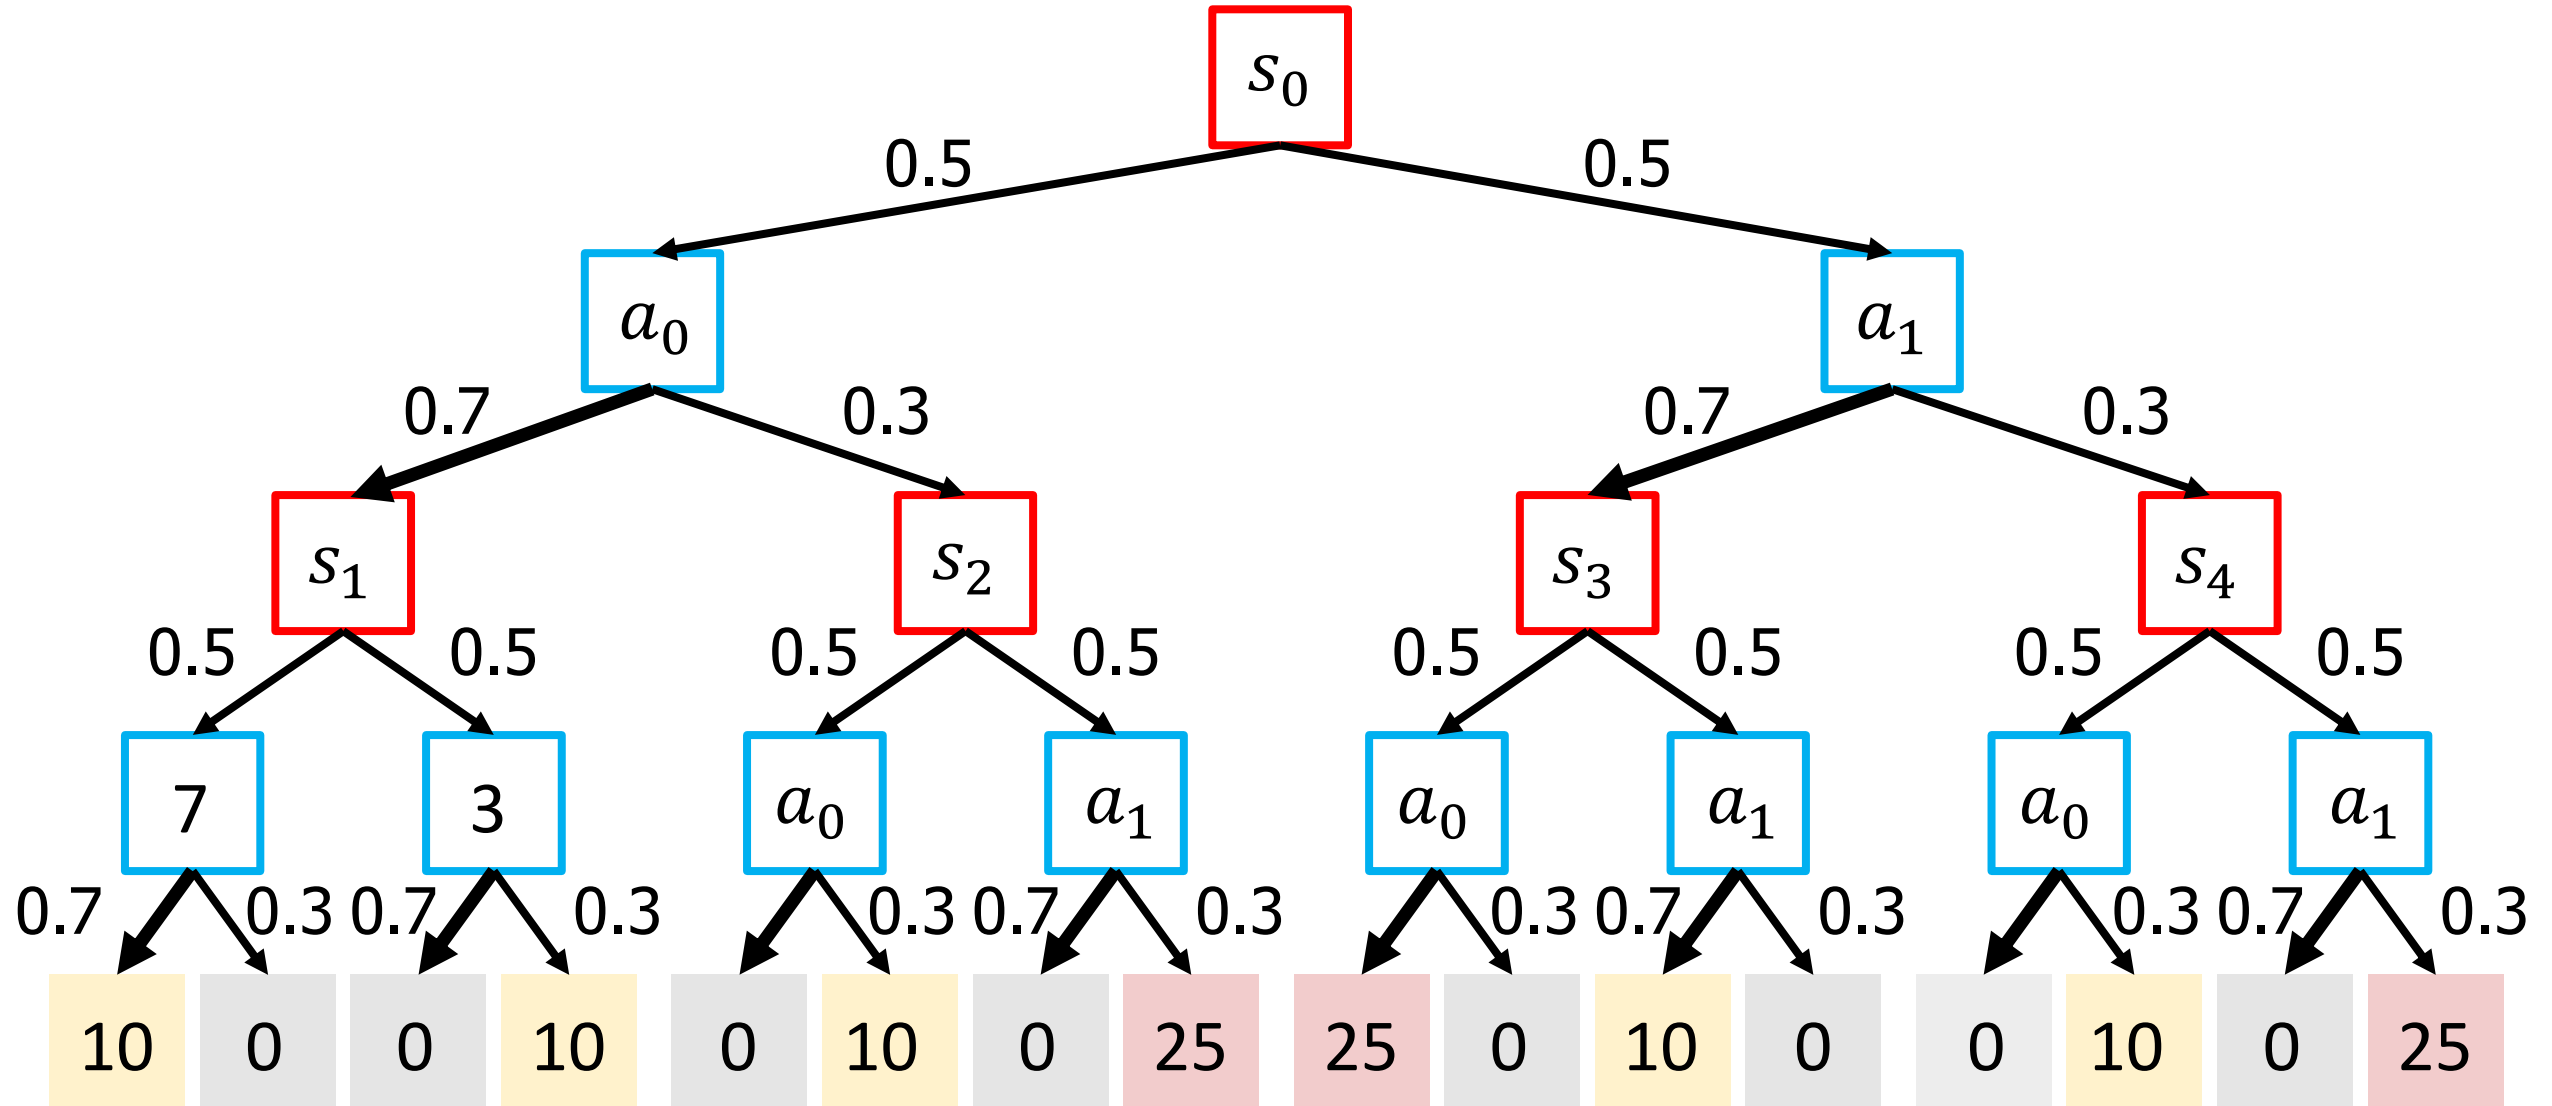
Compute the action-values at the second stage



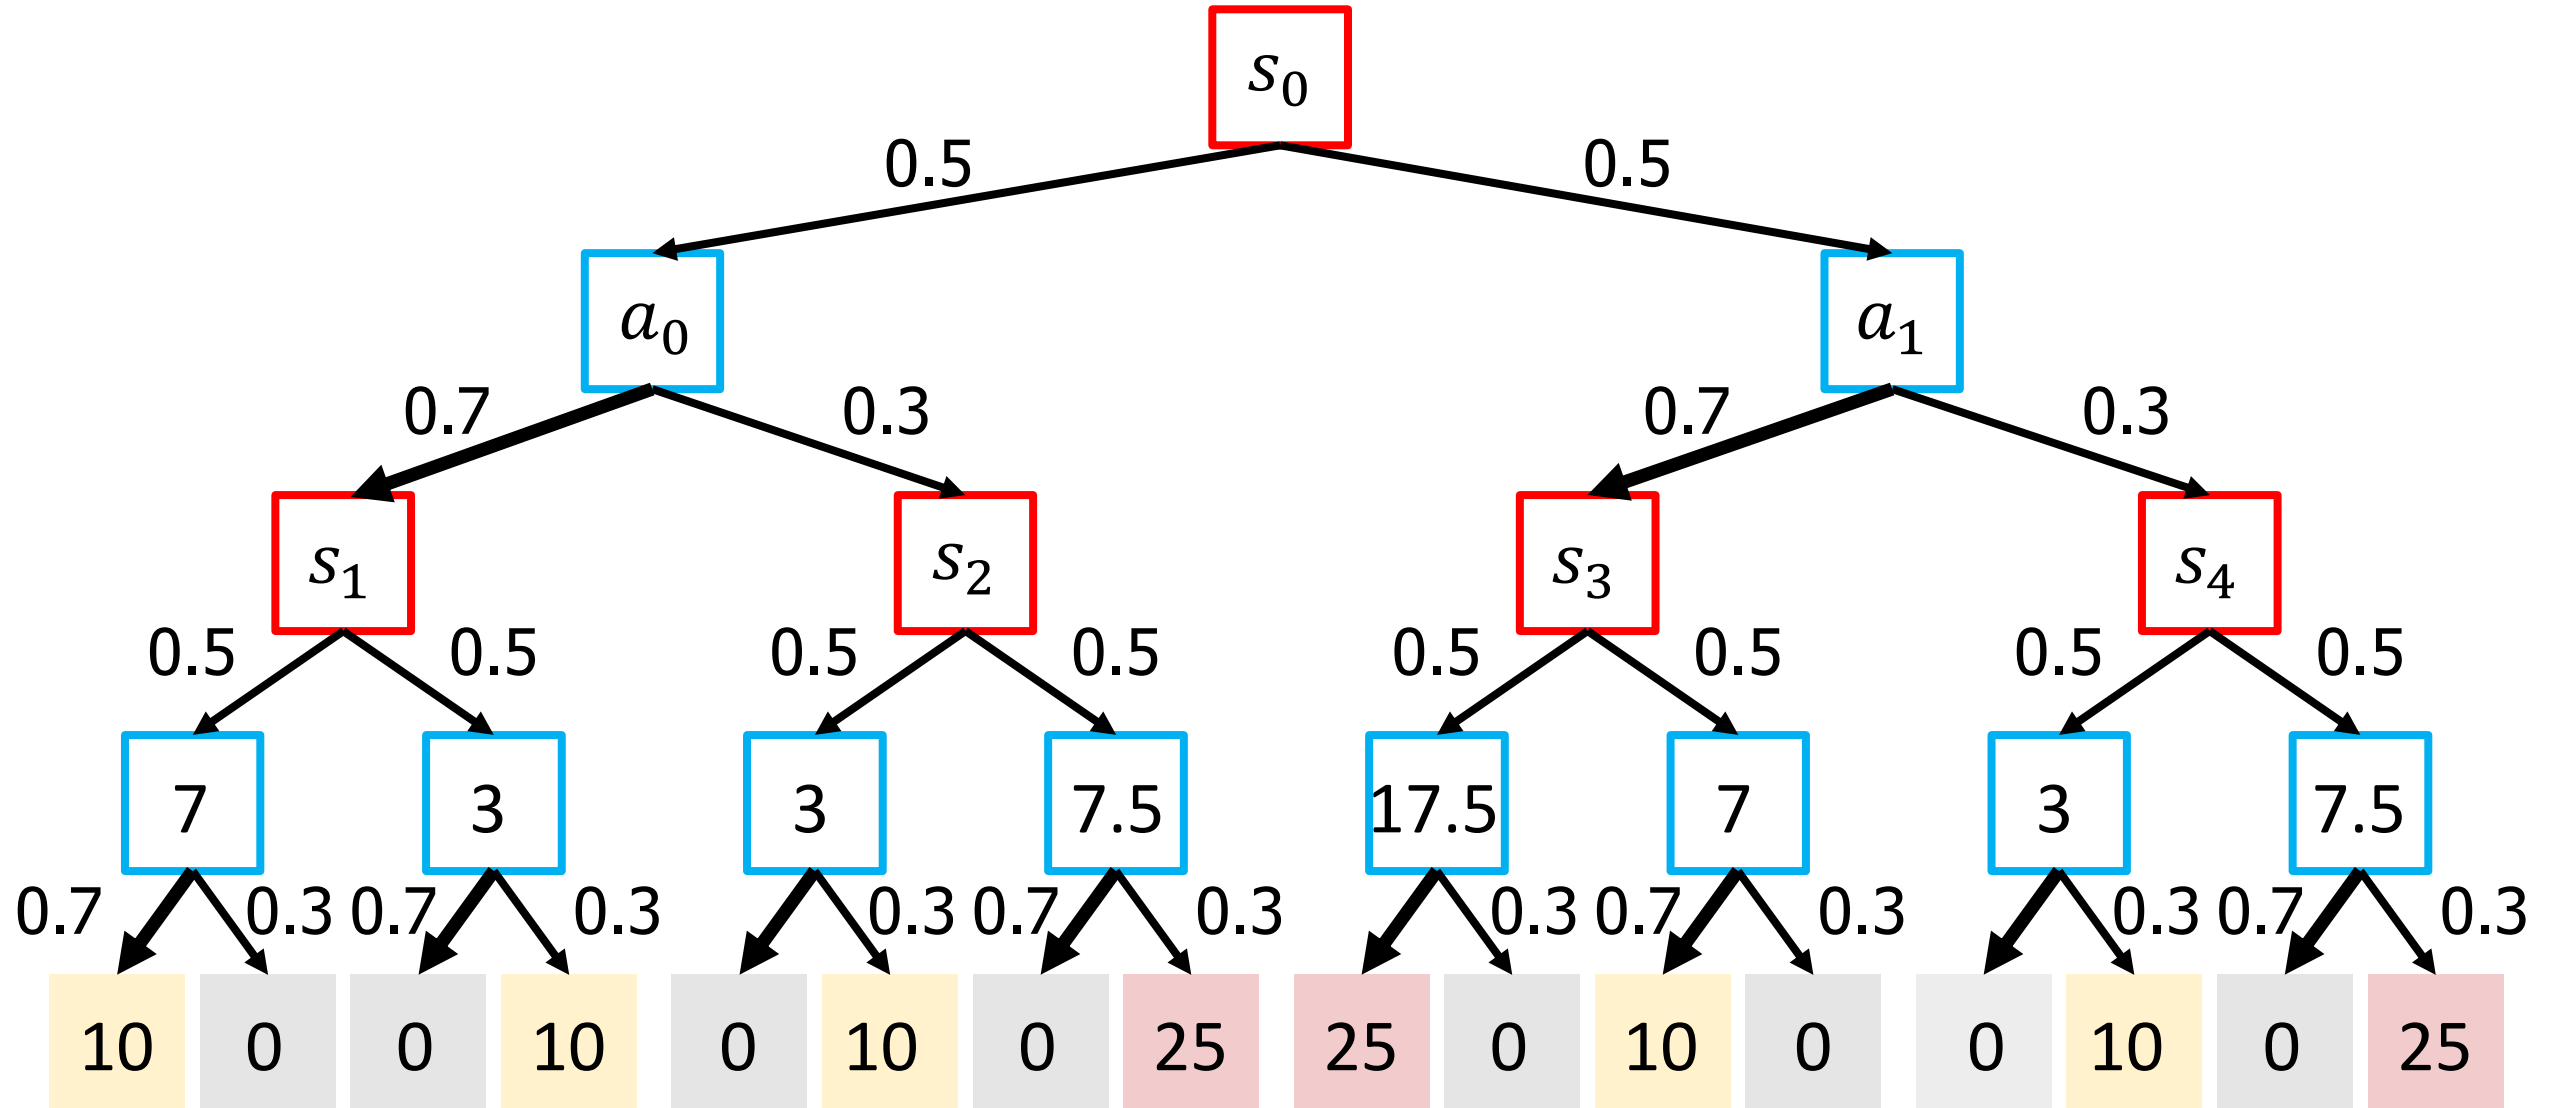
Compute the action-values at the second stage



Compute the action-values at the second stage



Compute the action-values at the second stage



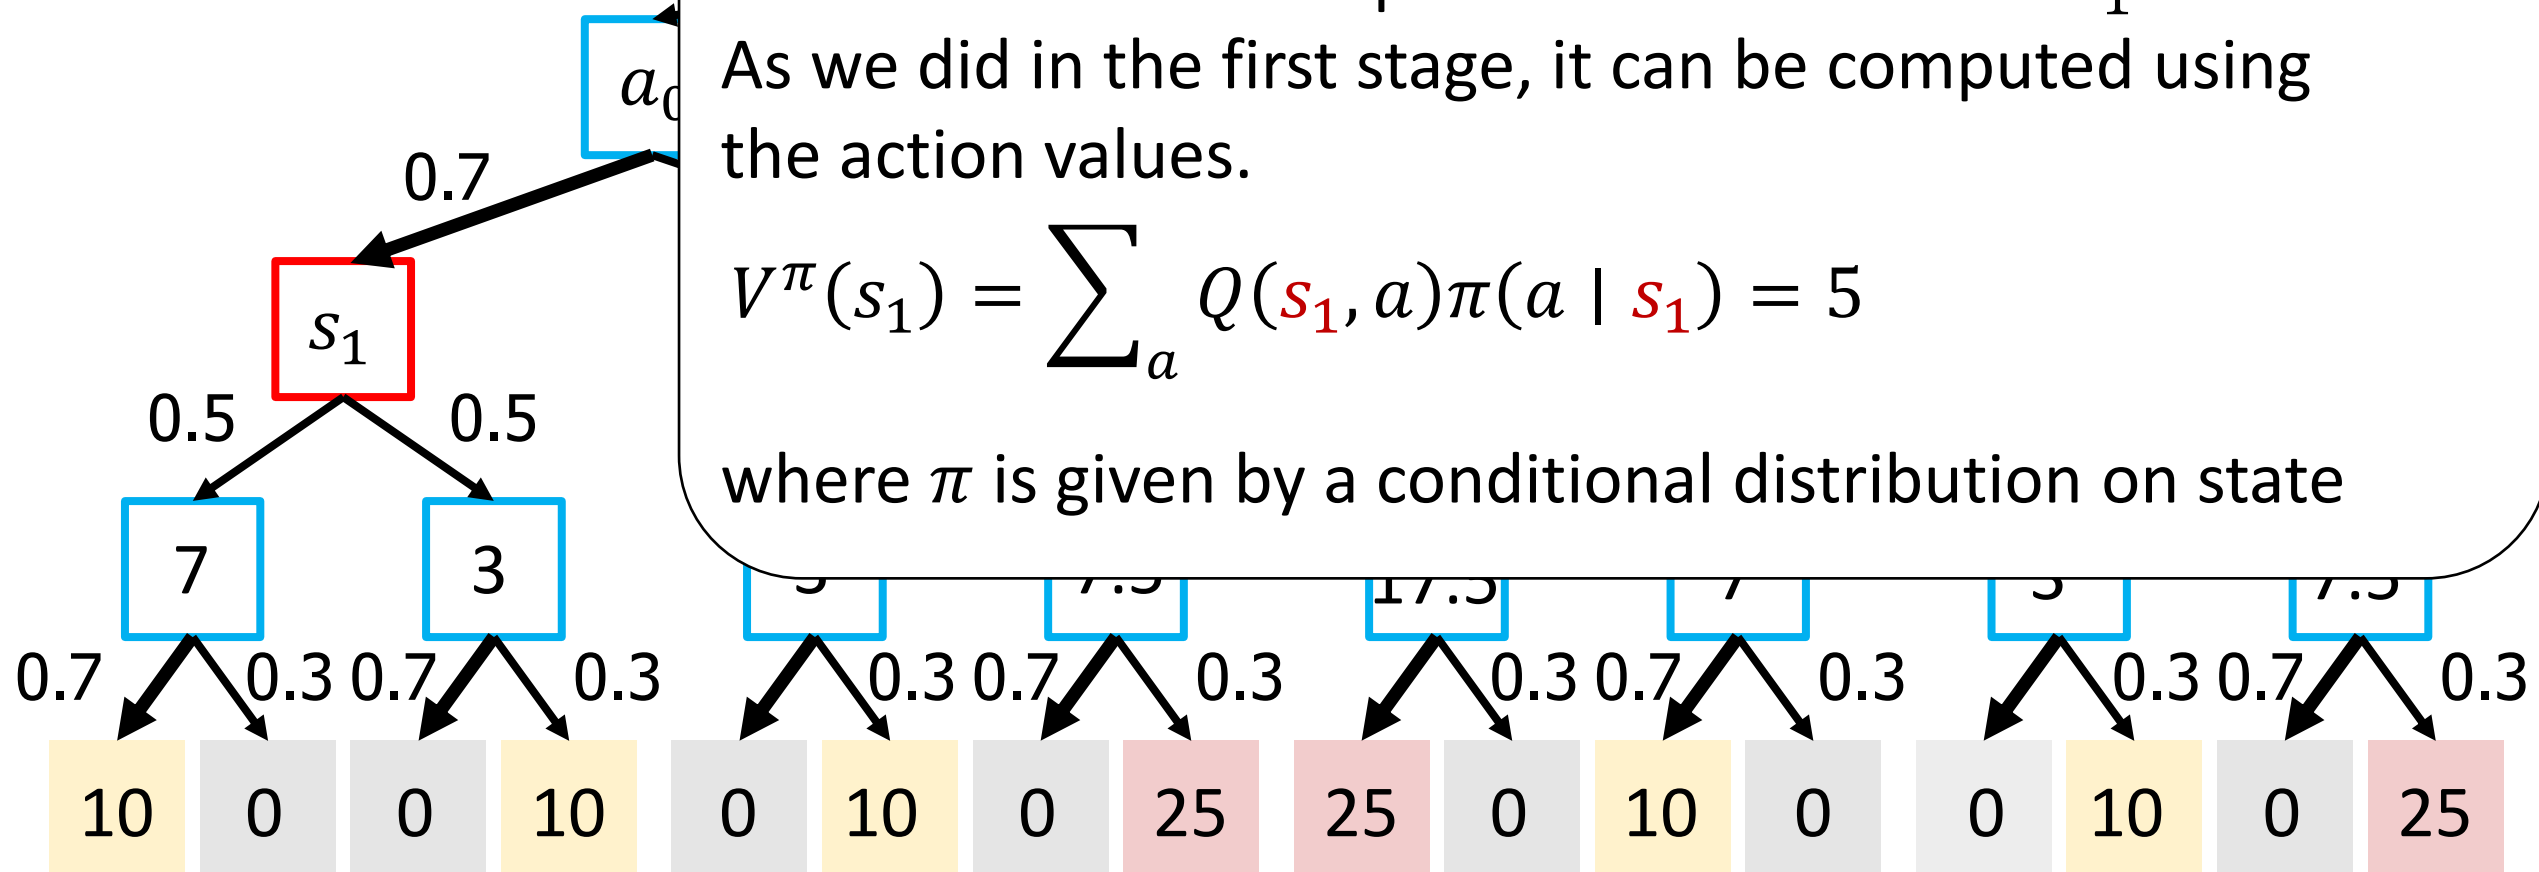
Compute the state-values at the second stage

s_1

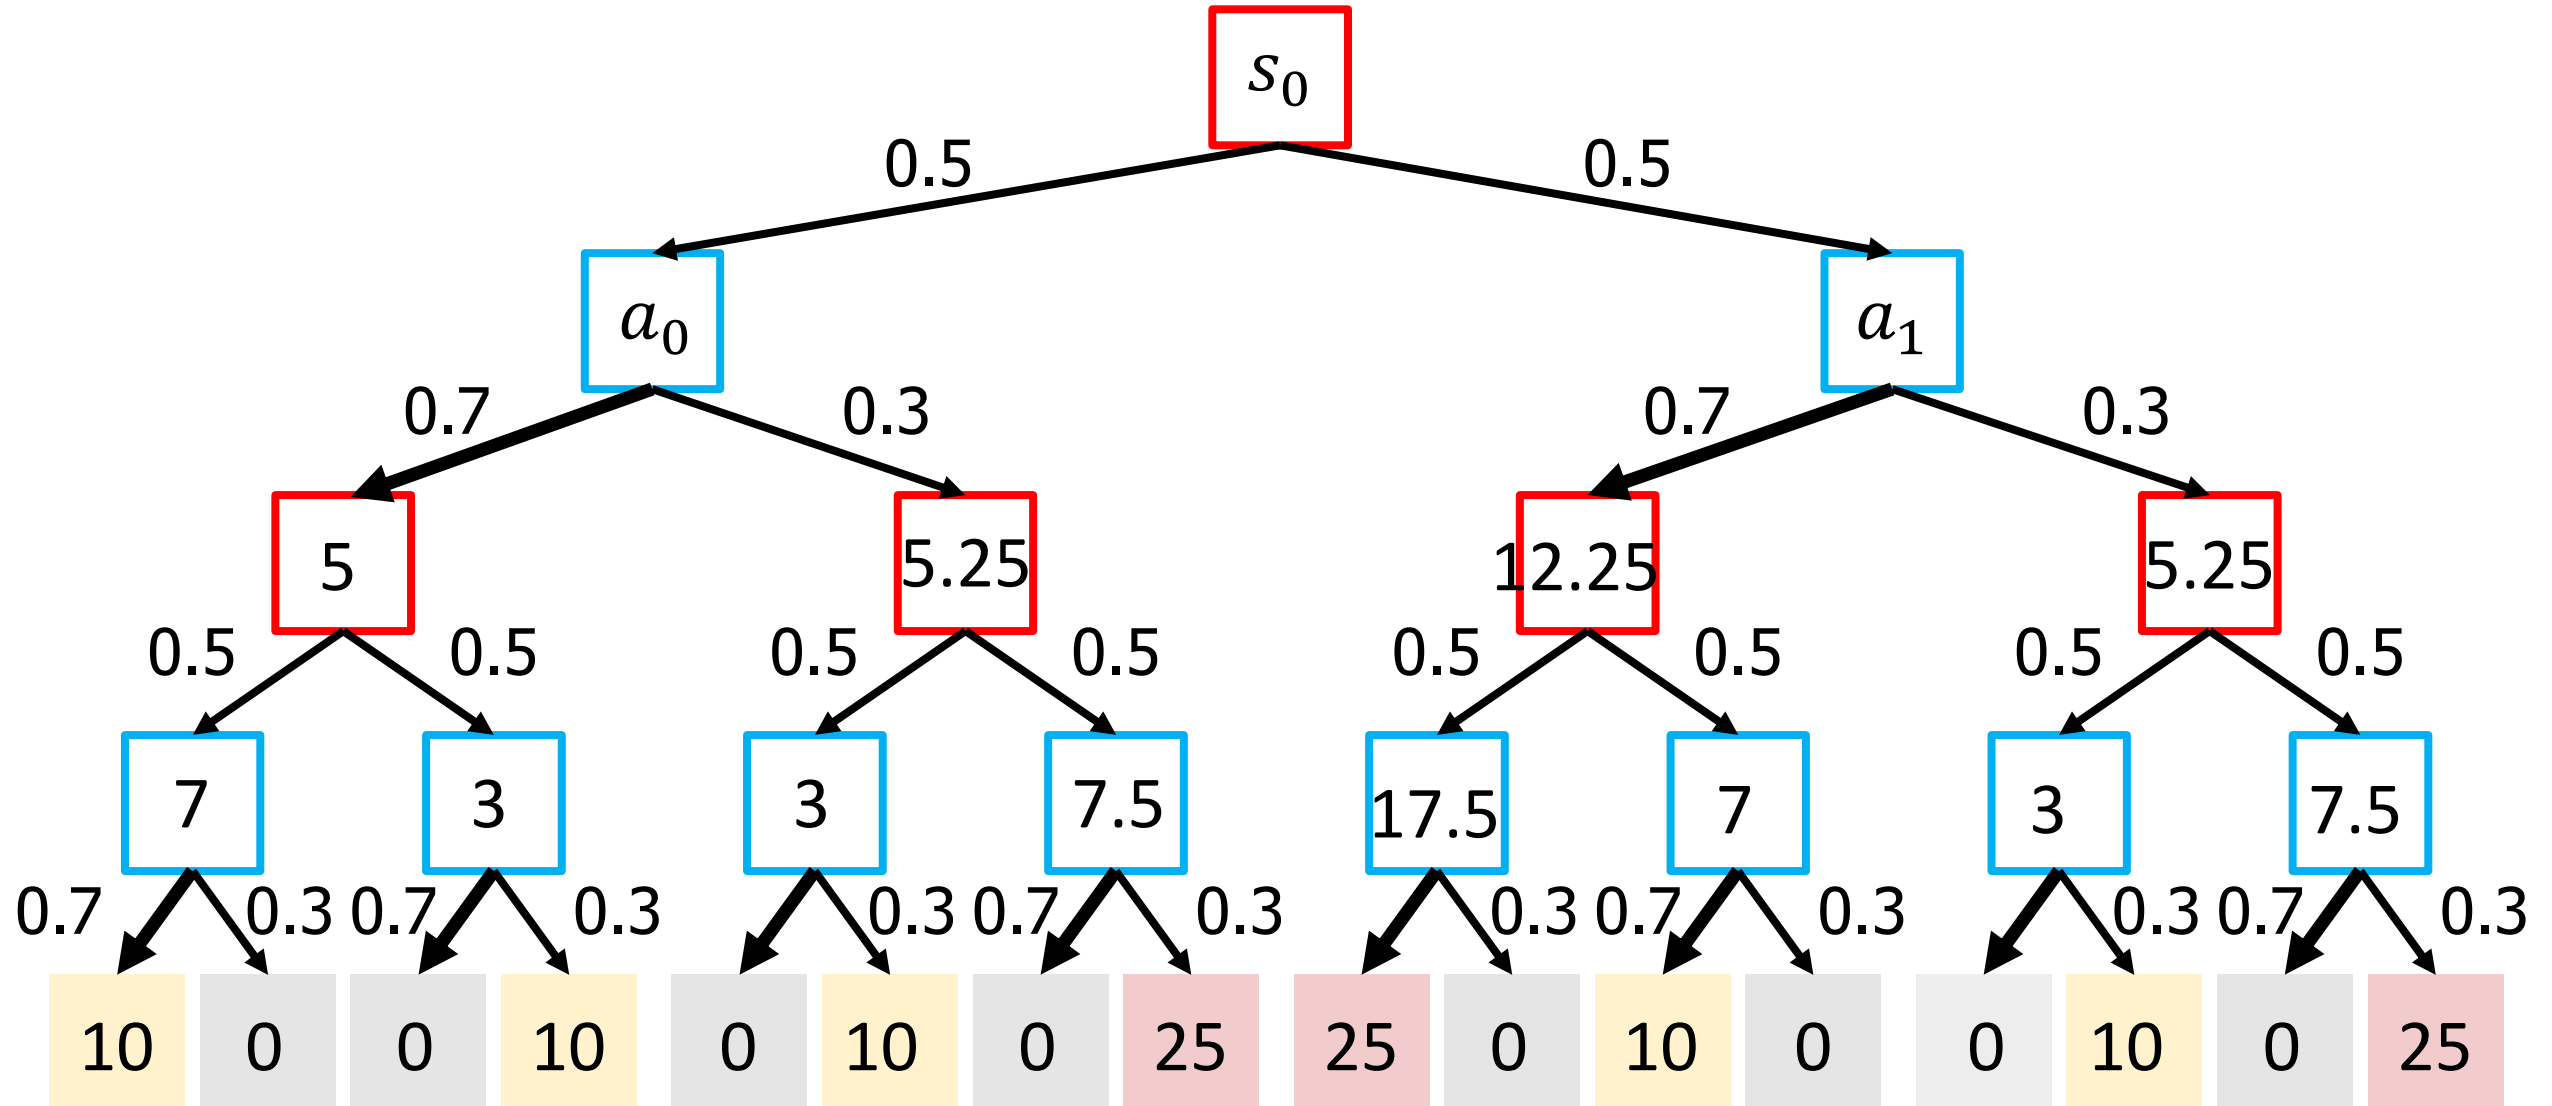
We would like to compute the state value at s_1 .
As we did in the first stage, it can be computed using the action values.

$$V^\pi(s_1) = \sum_a Q(s_1, a)\pi(a | s_1) = 5$$

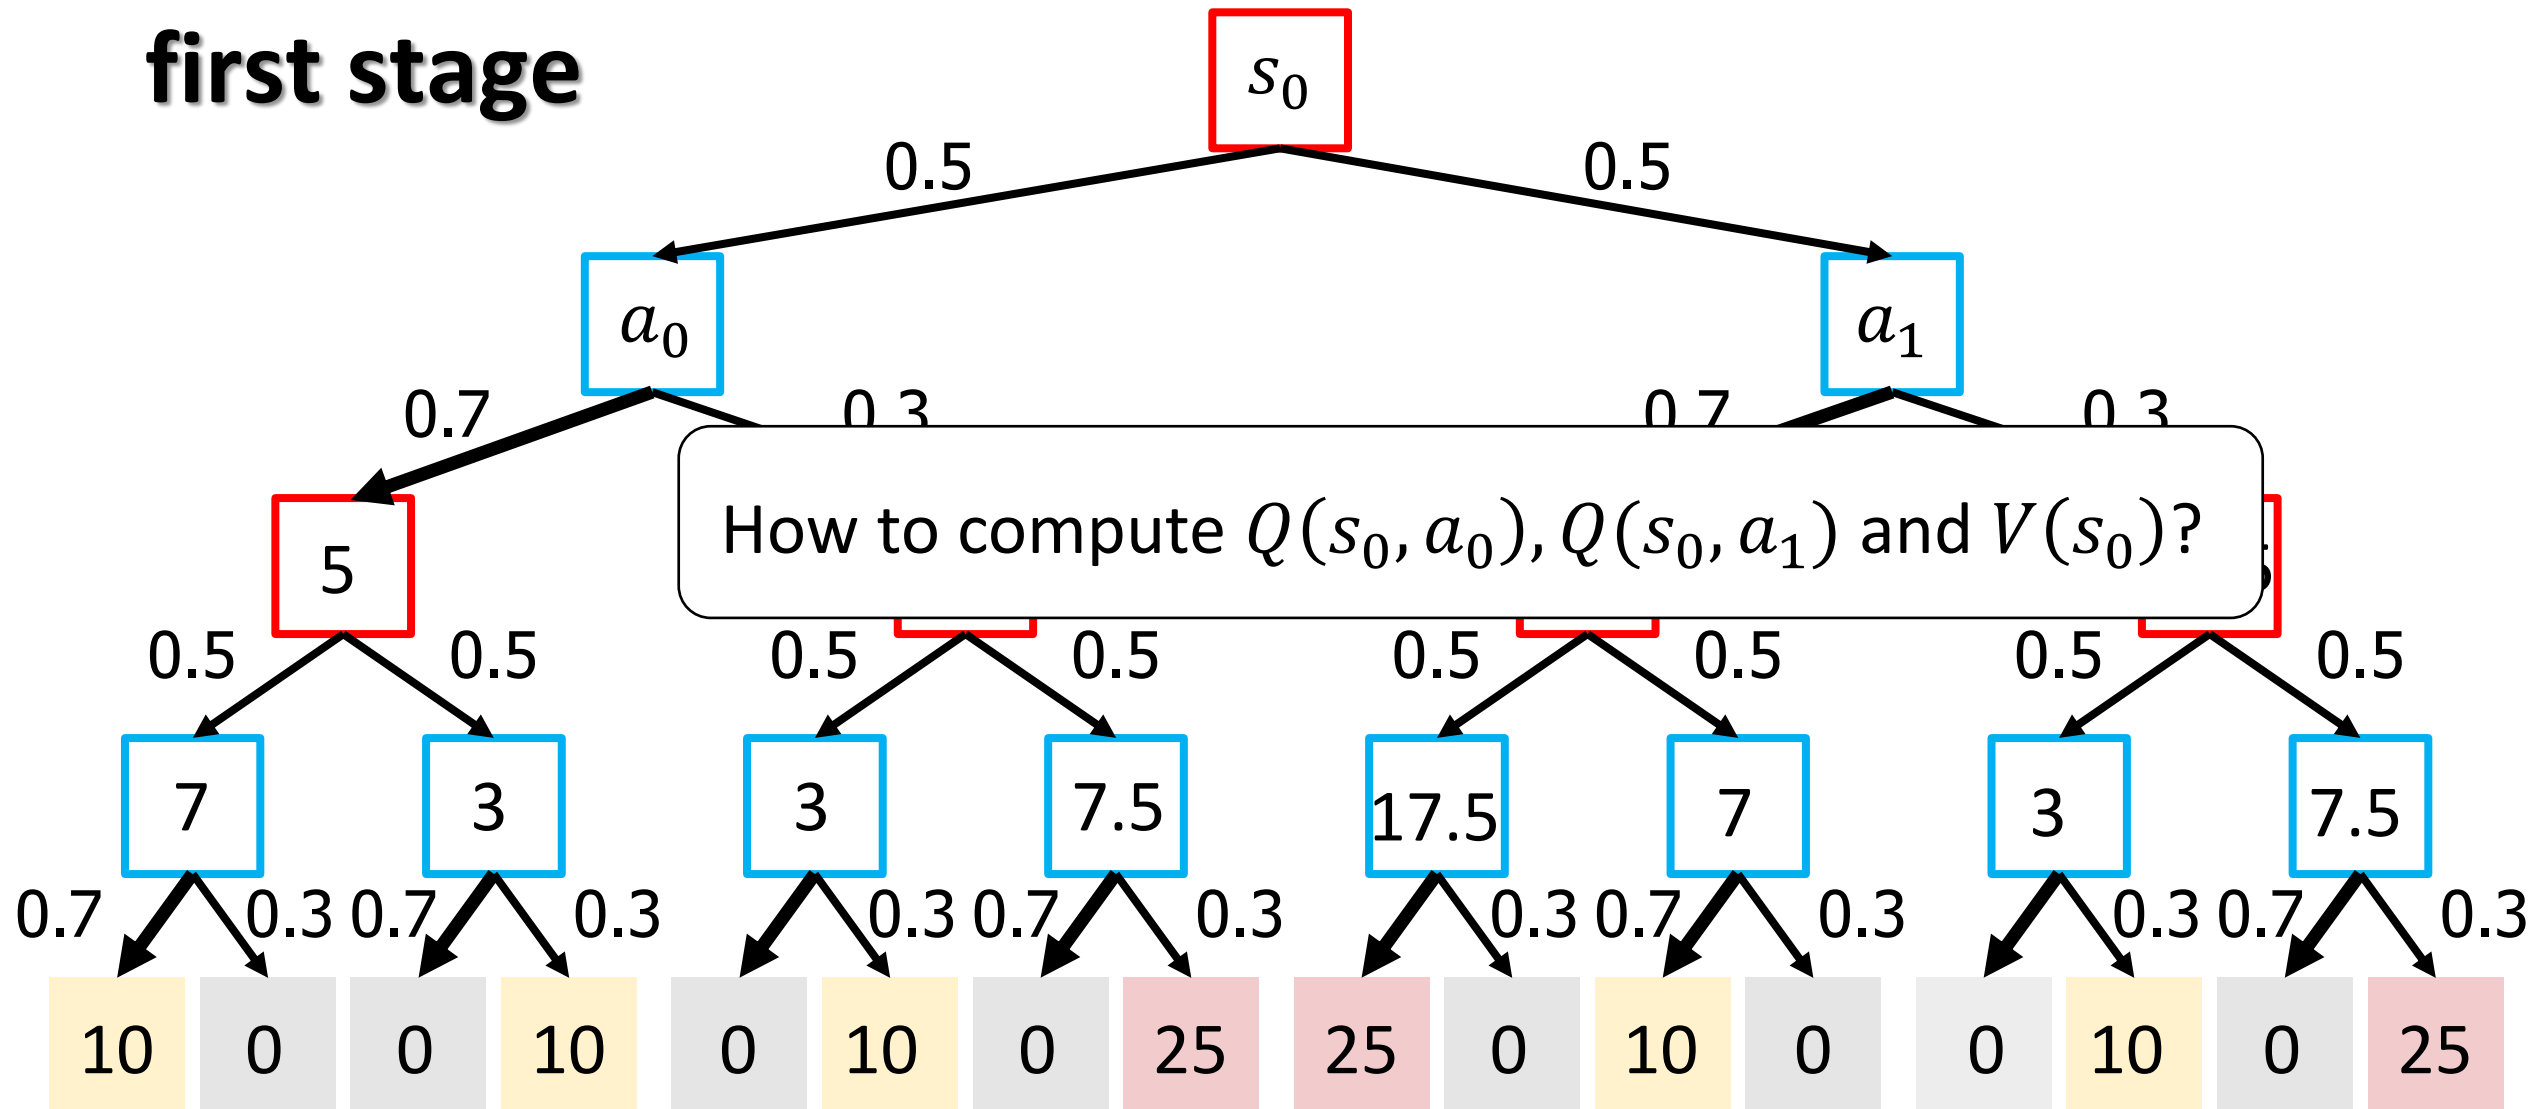
where π is given by a conditional distribution on state




Compute the state-values at the second stage



Compute the state- and action-values at the first stage



Compute the action-values at the first stage



• Possible transitions

- $s_1 \rightarrow a_0 \rightarrow 10$ with $0.7 \times 0.5 \times 0.7$
- $s_1 \rightarrow a_0 \rightarrow 0$ with $0.7 \times 0.5 \times 0.3$
- $s_1 \rightarrow a_1 \rightarrow 0$ with $0.7 \times 0.5 \times 0.7$
- $s_1 \rightarrow a_1 \rightarrow 10$ with $0.7 \times 0.5 \times 0.3$
- $s_2 \rightarrow a_0 \rightarrow 0$ with $0.3 \times 0.5 \times 0.7$
- $s_2 \rightarrow a_0 \rightarrow 10$ with $0.3 \times 0.5 \times 0.3$
- $s_2 \rightarrow a_1 \rightarrow 0$ with $0.3 \times 0.5 \times 0.7$
- $s_2 \rightarrow a_1 \rightarrow 25$ with $0.3 \times 0.5 \times 0.3$

Compute the action-values at the first stage

s_0


• Possible transitions

- $s_1 \rightarrow a_0 \rightarrow 10$ with $0.7 \times 0.5 \times 0.7$
- $s_1 \rightarrow a_0 \rightarrow 0$ with $0.7 \times 0.5 \times 0.3$
- $s_1 \rightarrow a_1 \rightarrow 0$ with $0.7 \times 0.5 \times 0.7$
- $s_1 \rightarrow a_1 \rightarrow 10$ with $0.7 \times 0.5 \times 0.3$
- $s_2 \rightarrow a_0 \rightarrow 0$ with $0.3 \times 0.5 \times 0.7$
- $s_2 \rightarrow a_0 \rightarrow 10$ with $0.3 \times 0.5 \times 0.3$
- $s_2 \rightarrow a_1 \rightarrow 0$ with $0.3 \times 0.5 \times 0.7$
- $s_2 \rightarrow a_1 \rightarrow 25$ with $0.3 \times 0.5 \times 0.3$

a_0

$$\begin{aligned} Q(s_0, a_0) &= 10 \times 0.7 \times 0.5 \times 0.7 \\ &\quad + 10 \times 0.7 \times 0.5 \times 0.3 \\ &\quad + 10 \times 0.3 \times 0.5 \times 0.3 \\ &\quad + 25 \times 0.3 \times 0.5 \times 0.3 \\ &= 5.075 \end{aligned}$$

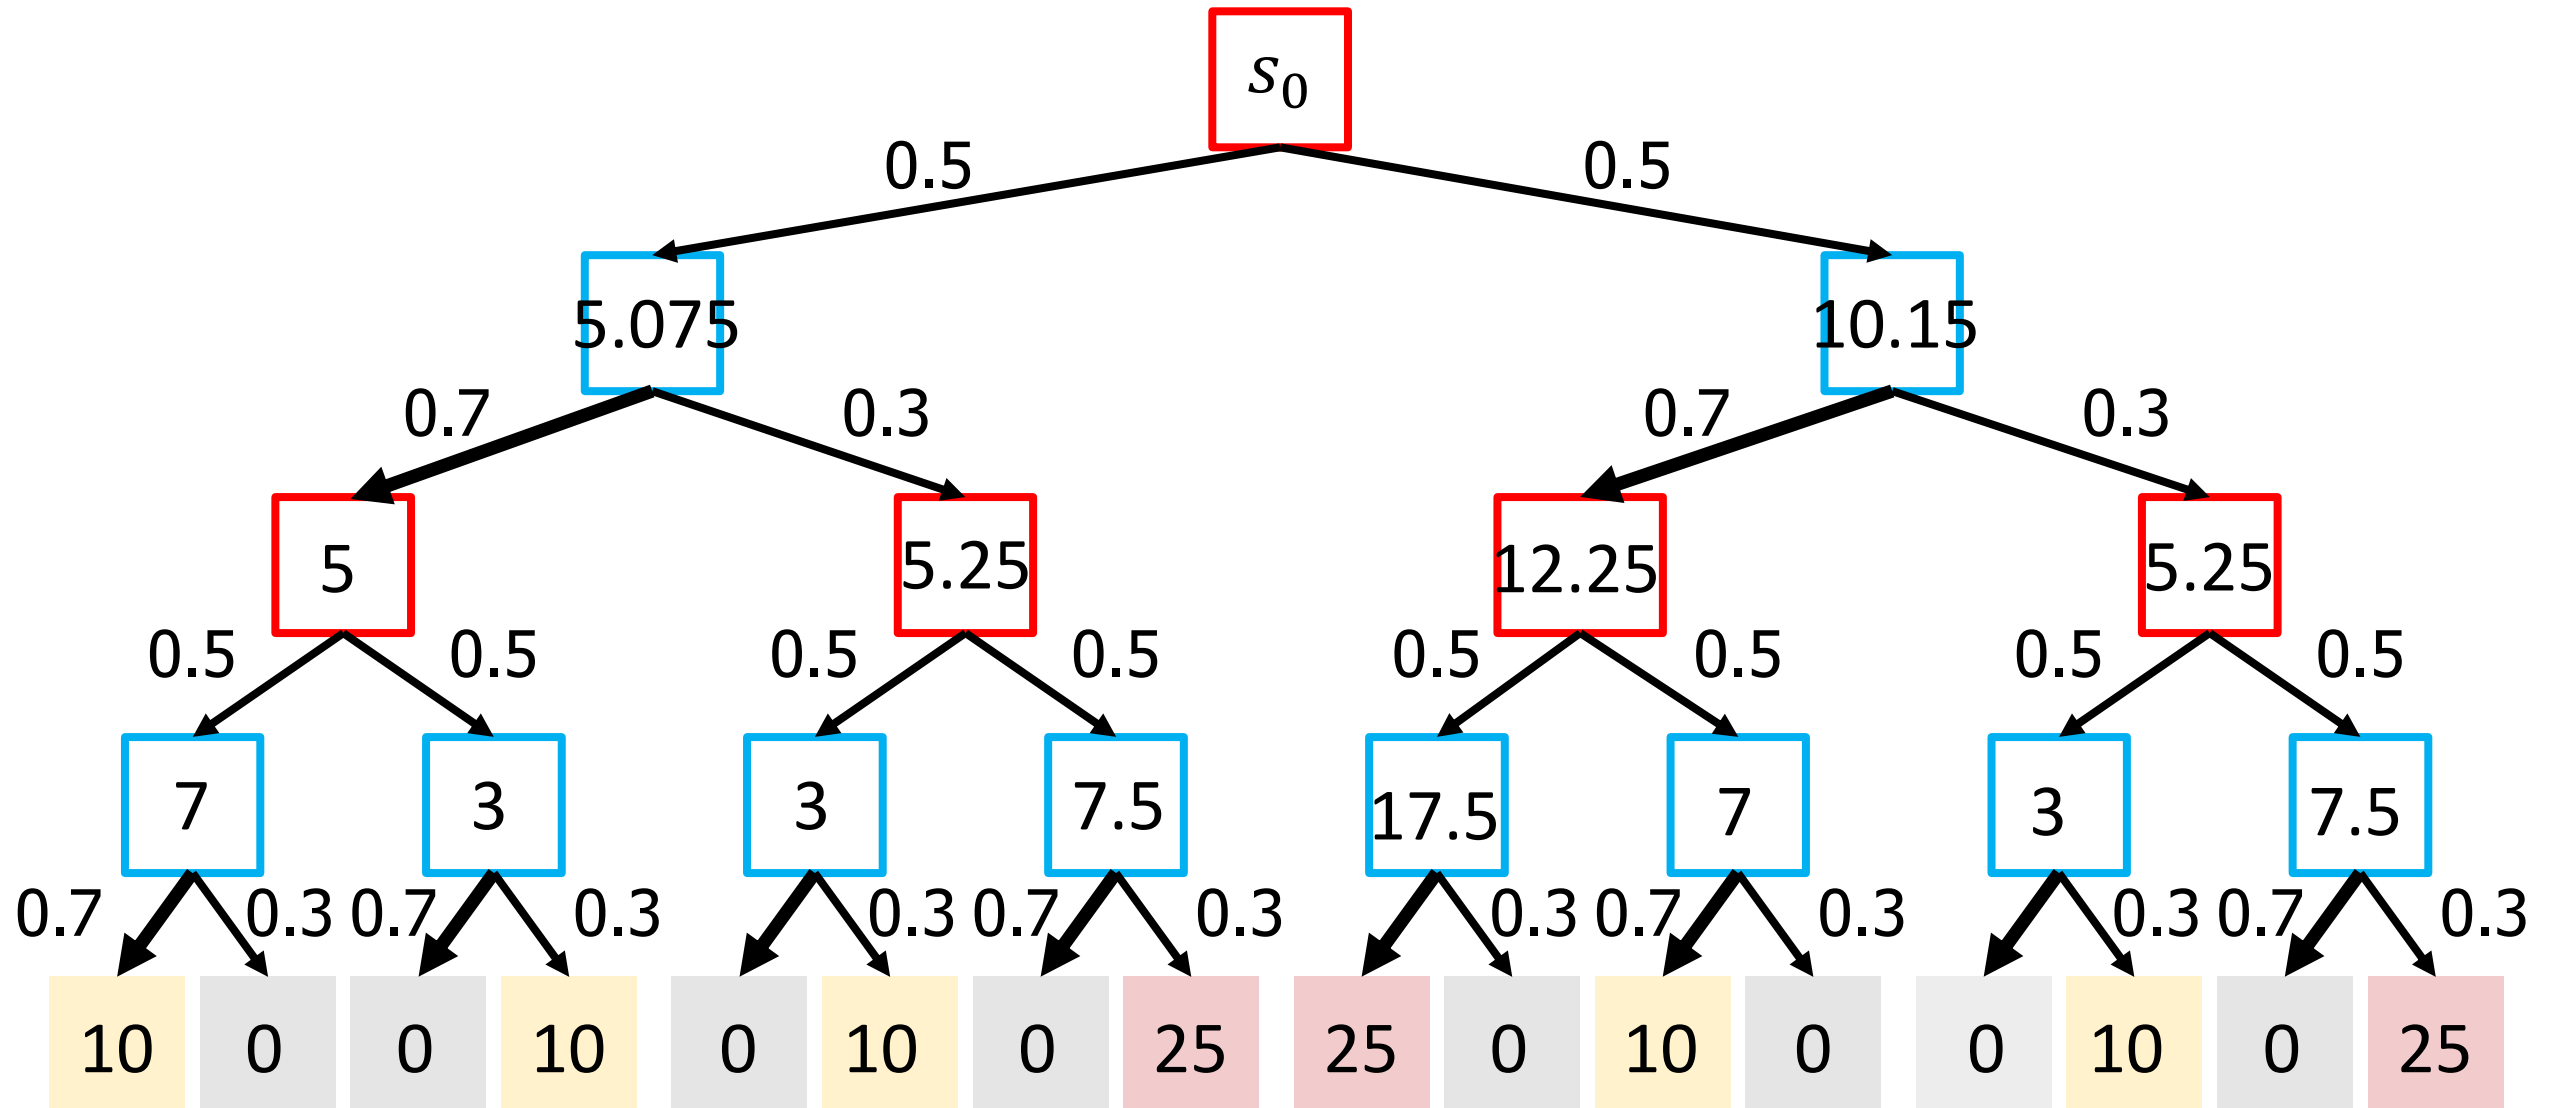
Compute the action-values at the first stage



• Possible transitions

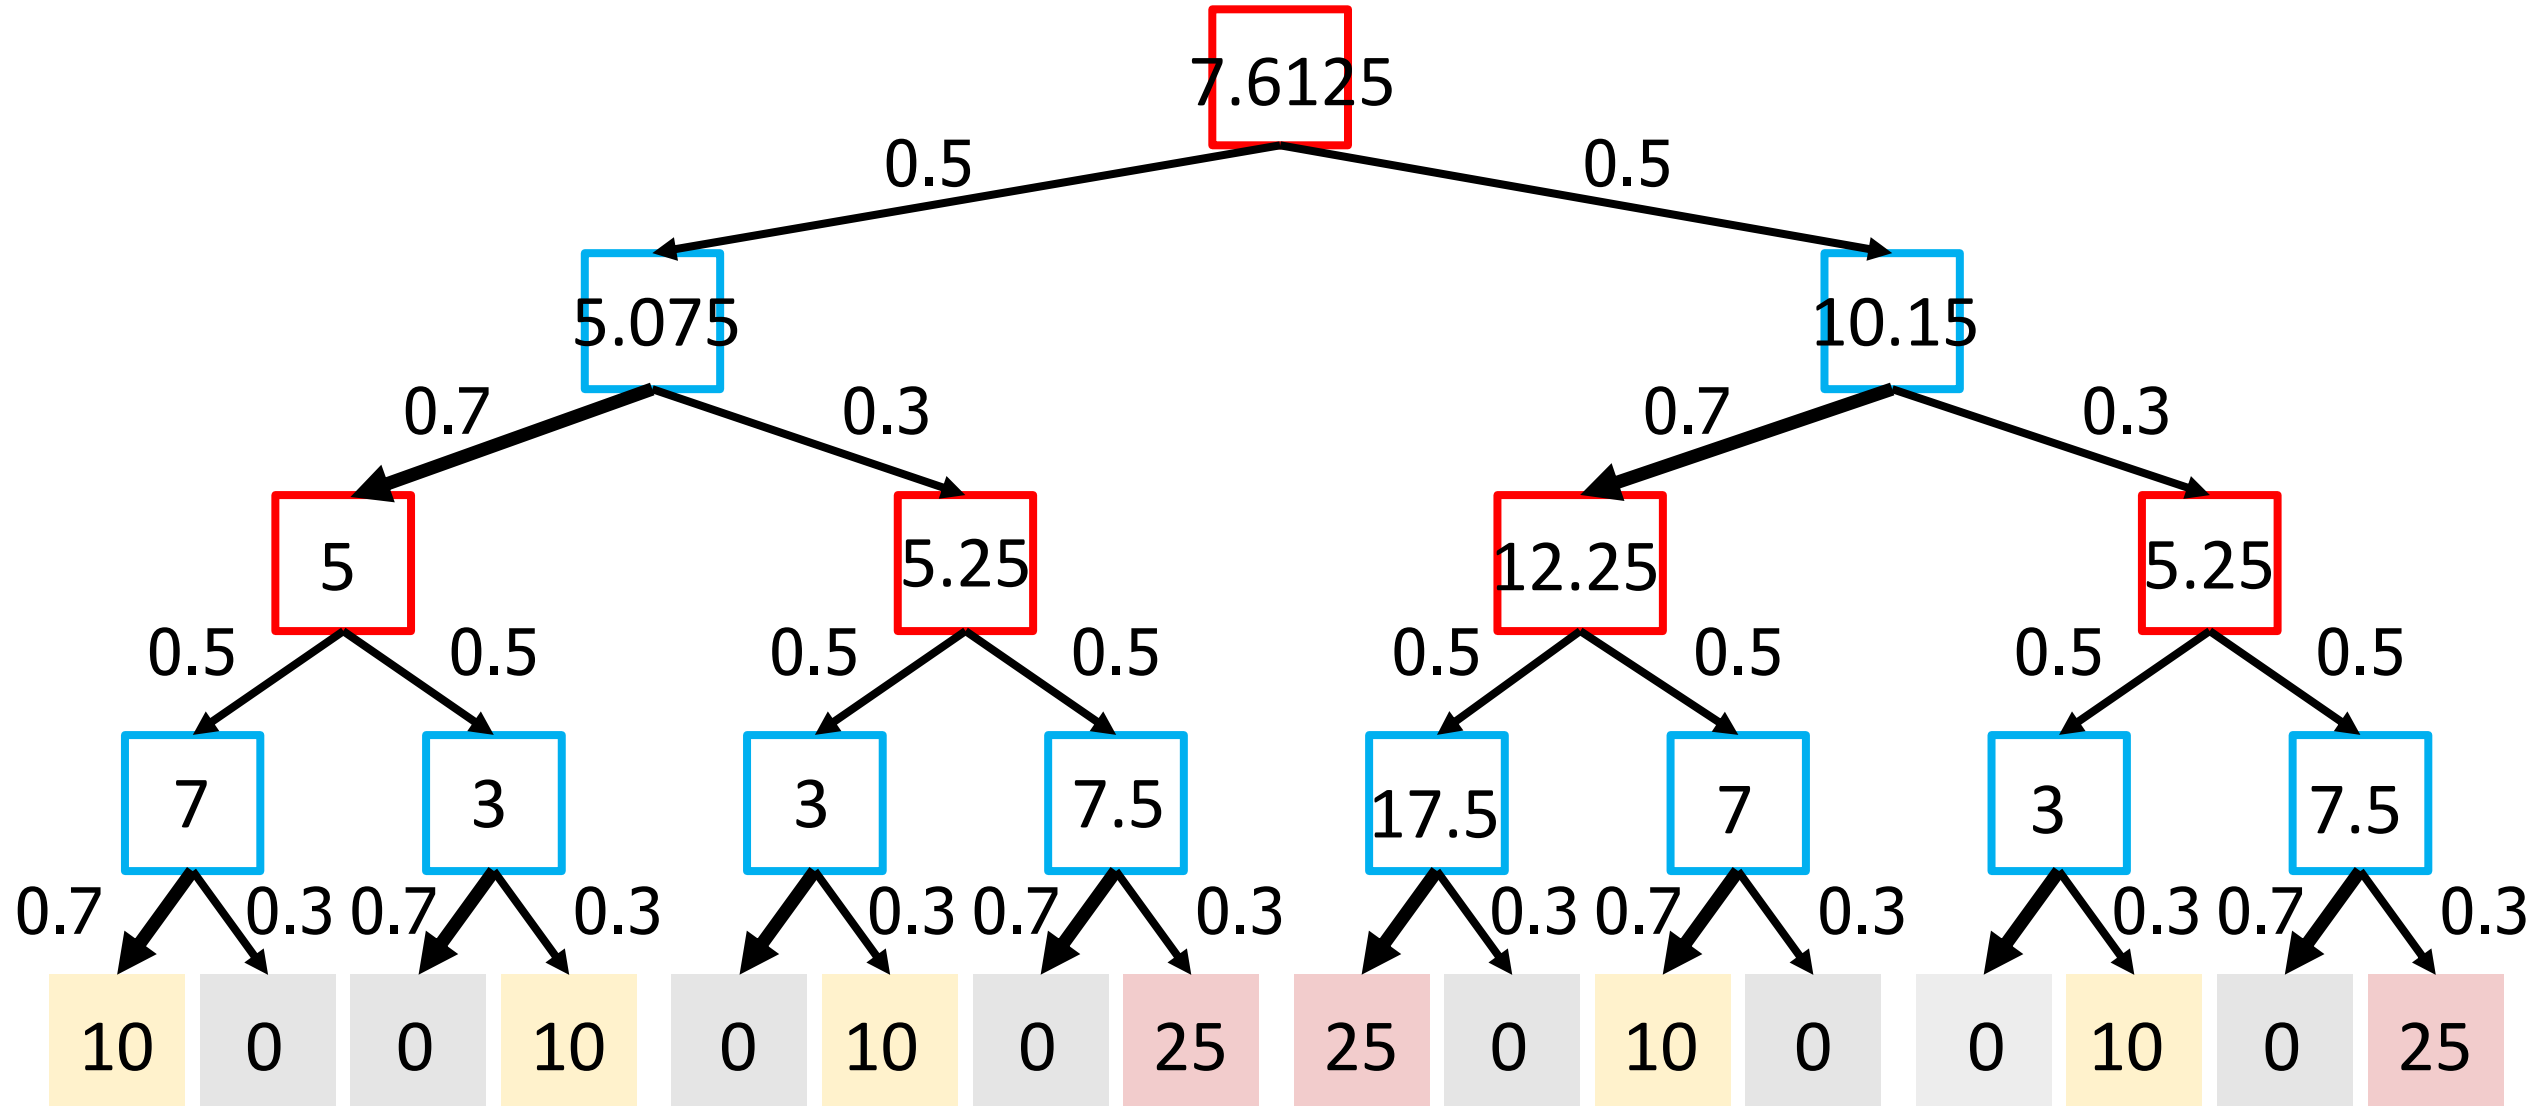
- $s_1 \rightarrow a_0 \rightarrow 10$ with $0.7 \times 0.5 \times 0.7$
- $s_1 \rightarrow a_0 \rightarrow 0$ with $0.7 \times 0.5 \times 0.3$
- $s_1 \rightarrow a_1 \rightarrow 0$ with $0.7 \times 0.5 \times 0.7$
- $s_1 \rightarrow a_1 \rightarrow 10$ with $0.7 \times 0.5 \times 0.3$
- $s_2 \rightarrow a_0 \rightarrow 0$ with $0.3 \times 0.5 \times 0.7$
- $s_2 \rightarrow a_0 \rightarrow 10$ with $0.3 \times 0.5 \times 0.3$
- $s_2 \rightarrow a_1 \rightarrow 0$ with $0.3 \times 0.5 \times 0.7$
- $s_2 \rightarrow a_1 \rightarrow 25$ with $0.3 \times 0.5 \times 0.3$

Compute the action-values at the first stage



Compute the state-value at the second stage


 [Open in Colab](#)





GENERAL FORM OF MARKOV DECISION PROCESS


Extension to Sequential Decision Making Problems

- A learning agent is not in front of gambling machines
- A current position is considered as a state
 - The state is a sufficient statistics to describe the dynamics of the environment
 - Bandits are MDPs with one state
- The action of the agent influences the environment, causing a state transition




Extension to Sequential Decision Making Problems

- s : state, 
- a : action, 
- State set: $\mathcal{S} = \{s_1, s_2, \dots, s_{|\mathcal{S}|}\}$
- Action set: $\mathcal{A} = \{a_1, a_2, \dots, a_{|\mathcal{A}|}\}$



Extension to Sequential Decision Making Problems

- $p_0(s)$: initial state distribution
- $p(s', r \mid s, a)$: stochastic environmental dynamics
probability of occurring state s' and reward r
when executing action a at state s
- $\pi(a \mid s)$: stochastic policy
probability to select action a at state s
- $\pi(s) \in \mathcal{A}$: deterministic policy
action at state s



Markov property

- Consider a sequence of states, action, and rewards

$$s_0, a_0, r_1, s_1, a_1, r_2, \dots, s_{t-1}, a_{t-1}, r_t, s_t, a_t, s_{t+1}, r_{t+1}$$

where the subscript t represents the time index

- How can we model the next state s_{t+1} and the reward r_{t+1} ?
The most general way is to introduce a conditional probability distribution

$$\Pr(s_{t+1} = s', r_{t+1} = r \mid s_0, a_0, r_1, \dots, s_{t-1}, a_{t-1}, r_t, s_t, a_t)$$

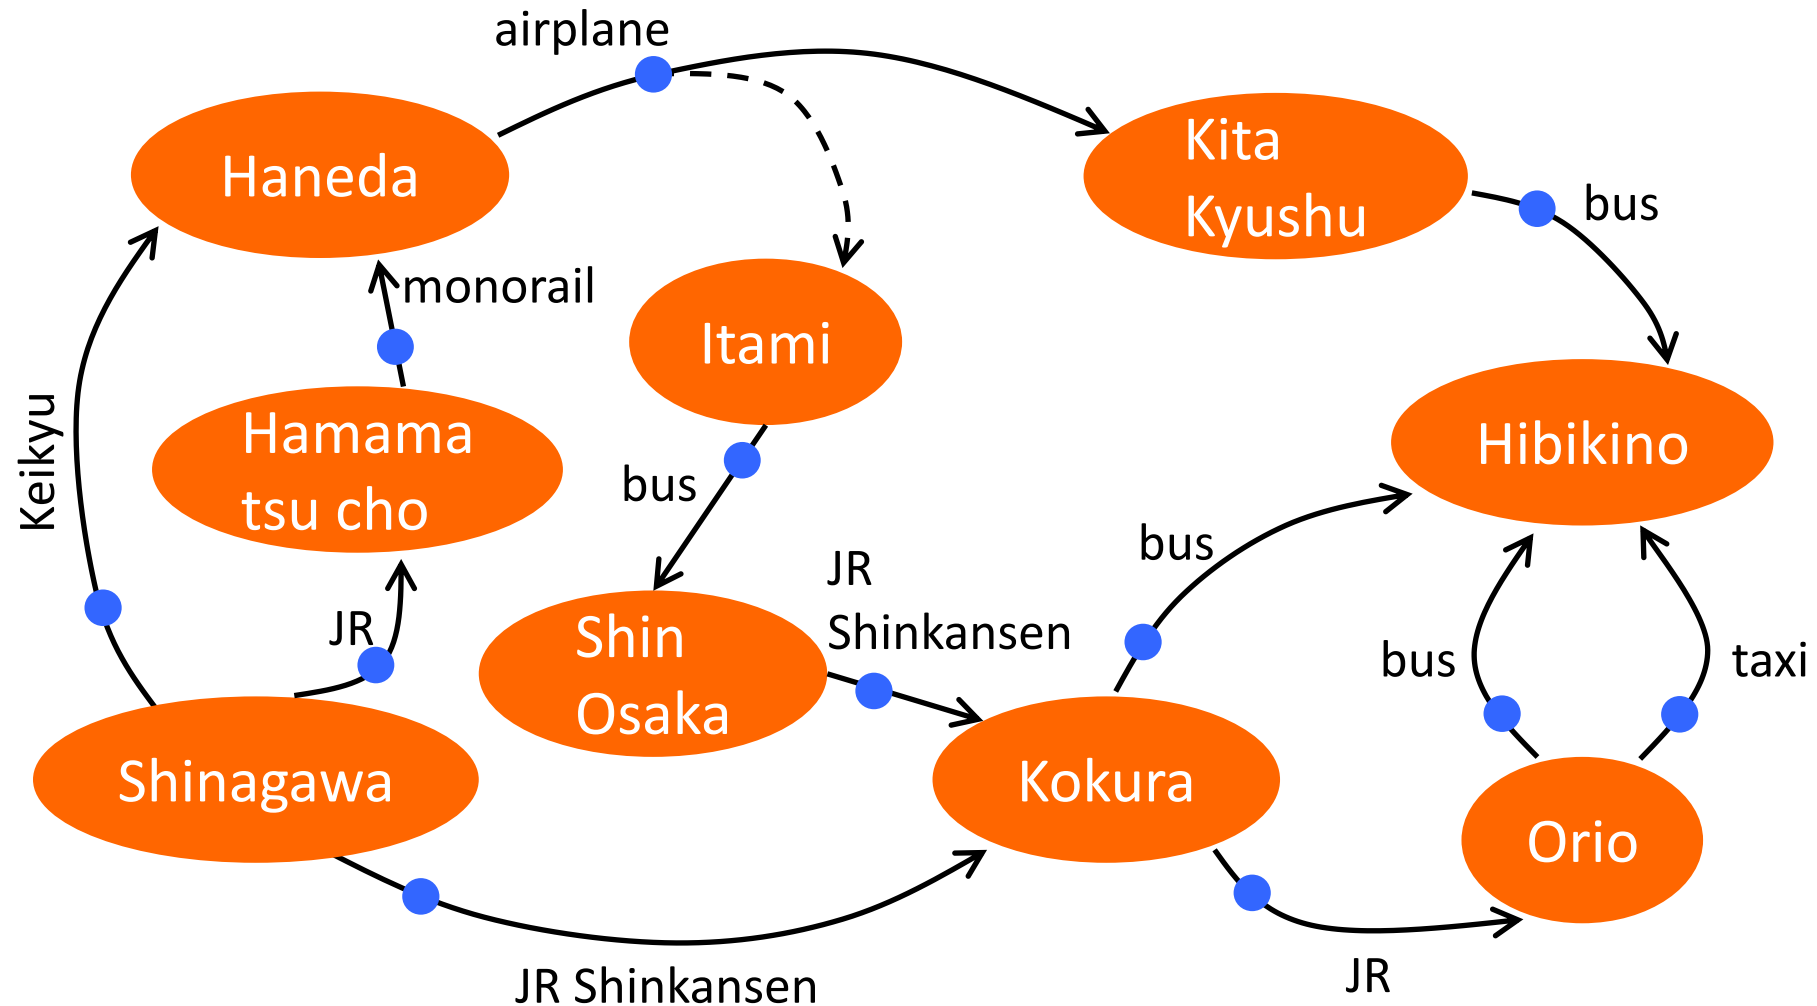
Markov property

- It is too complicated! We usually **assume** that the environment satisfies the following Markov property:

$$p(s', r \mid s, a) = \Pr(S_{t+1} = s', R_{t+1} = r \mid S_t = s, A_t = a)$$

- We have already utilized the Markovian assumption in the previous two-stage task

Example: From Shinagawa to Hibikino Campus



States and Actions in the Example

- $\mathcal{S} = \{\text{Shinagawa, Hamamatsu cho, Haneda, Itami, Shin Osaka, Kitakyushu, Kokura, Orio, Hibikino}\}$

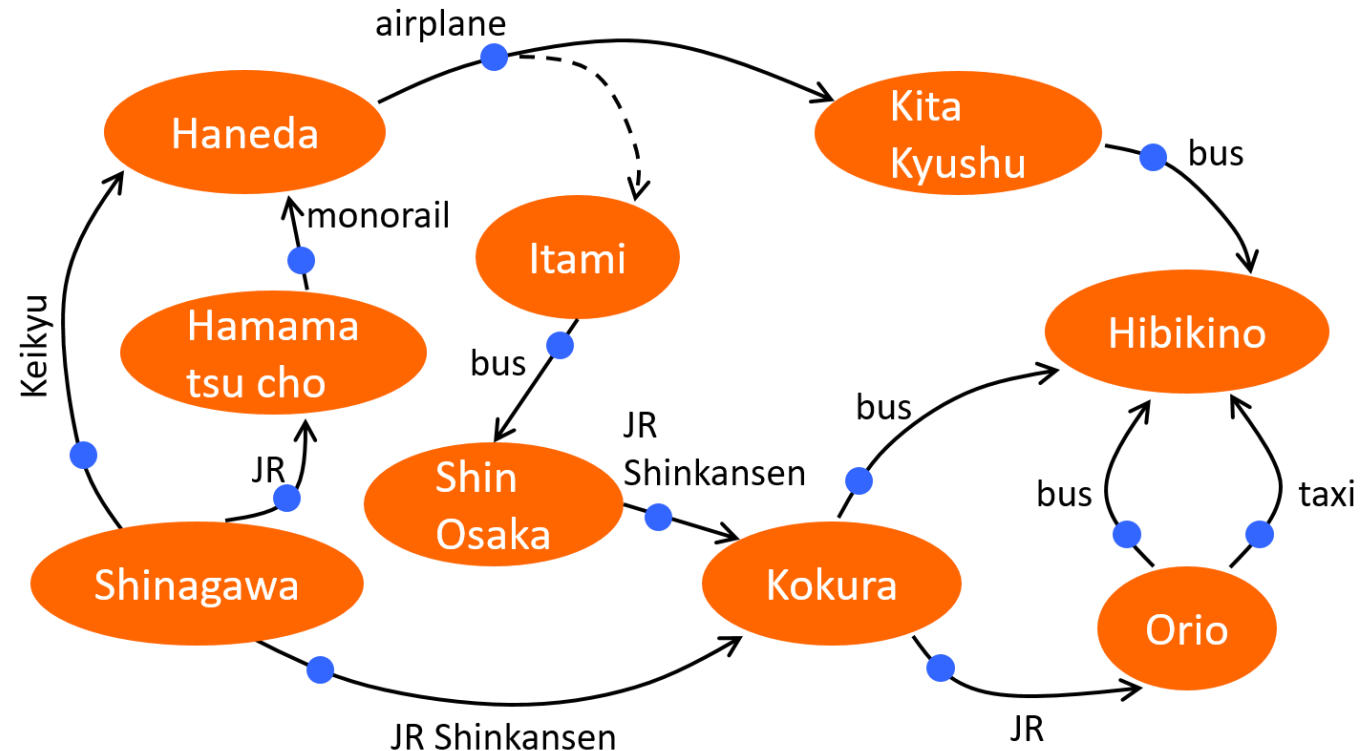
- Action set is state dependent in this case

- $\mathcal{A}(\text{Shinagawa}) = \{\text{Keikyu, JR, JR Shinkansen}\}$

- $\mathcal{A}(\text{Kokura}) = \{\text{bus, JR}\}$

- $\mathcal{A}(\text{Haneda}) = \{\text{Airplane}\}$

- \vdots



State Representation is Important

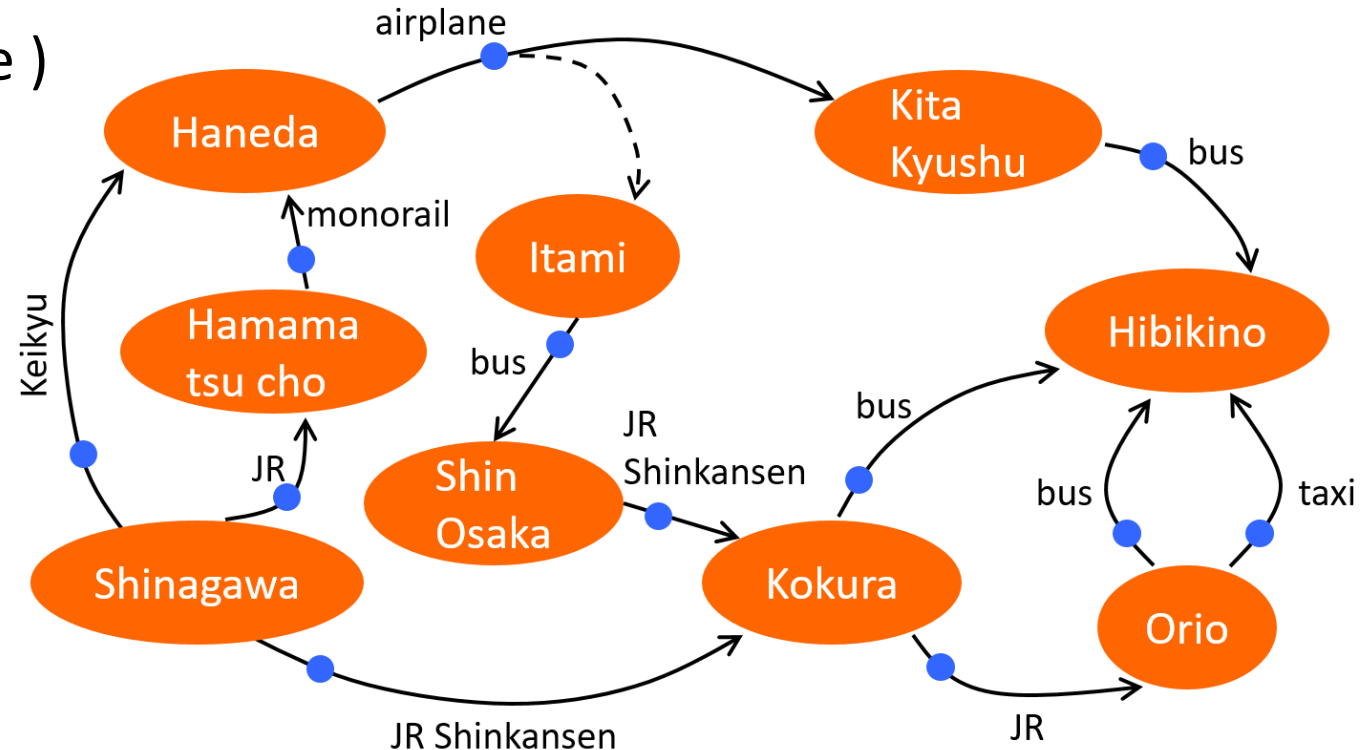
- Markovian property

- $\Pr(\text{Kita Kyushu} | \text{Haneda, airplane, Hamamatsu-cho, monorail, Shinagawa, JR})$

- $= \Pr(\text{Kita Kyushu} | \text{Haneda, airplane, Shinagawa, Keikyu})$


- $= \Pr(\text{Kita Kyushu} | \text{Haneda, airplane})$

independent on the route



State Representation is Important

- The task of the robot is to move from the start position to the goal (red cylinder)
- The robot cannot see the red cylinder at the corner
- If a state is given by a position and an orientation of the robot, simple RL algorithms cannot achieve this task because the state representation does not satisfy Markov property
- Memory is needed



J. Wang et al. (2020). [Modular deep reinforcement learning from reward and punishment for robot navigation](#). Neural Networks, 135: 115-126.

Components

- State transition probability: $p(s' | s, a) = \sum_r p(s', r | s, a)$
- Expected reward for state-action pairs:

$$r(s, a) = \sum_r r p(r | s, a) \quad \text{where} \quad p(r | s, a) = \sum_{s'} p(s', r | s, a)$$

- Expected reward for state-action-next-state triples


$$r(s, a, s') = \sum_r r p(r | s, a, s') \quad \text{where} \quad p(r | s, a, s') = \frac{p(s', r | s, a)}{p(s' | s, a)}$$

Return

- The return G_t is defined as the sum of discounted rewards from time-step t

$$\begin{aligned} G_t &= r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \\ &= \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \end{aligned}$$

- G_t is a random variable
- $\gamma \in [0, 1)$ is the discount rate
 - γ close to 0: myopic evaluation
 - γ close to 1: far-sighted evaluation




Why Discount?

- The learned behavior is affected by the choice of γ



The robot does not move towards the battery



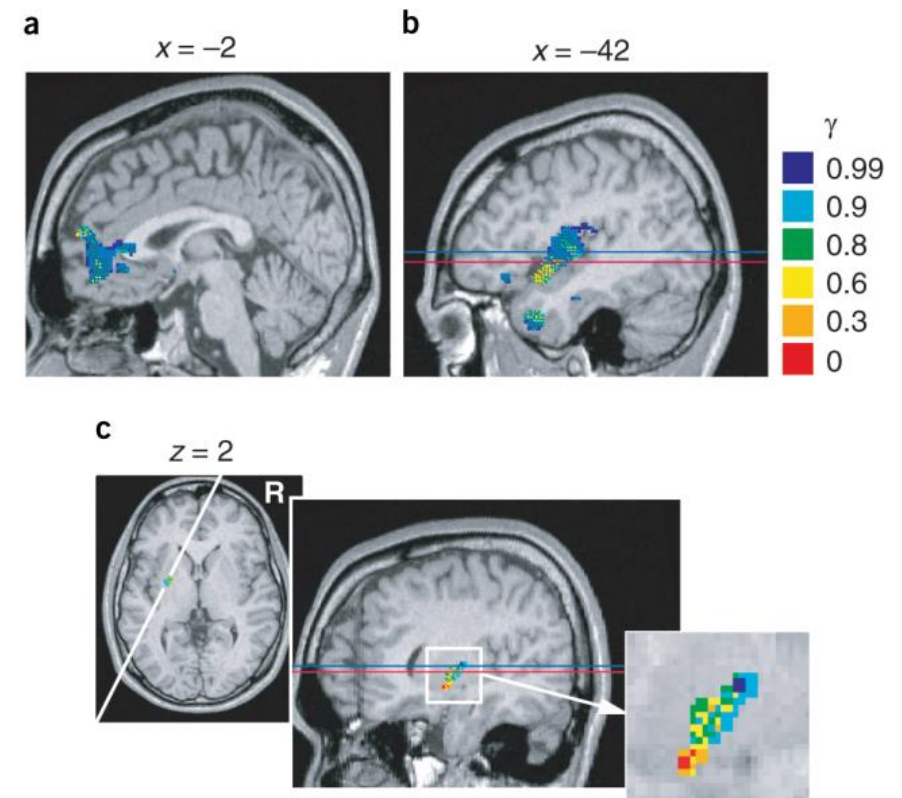
The robot tries to catch the battery

- If the reward is bounded, the return is also bounded:

$$|R| \leq R_{\max} \quad \longrightarrow \quad |G| \leq \frac{R_{\max}}{1 - \gamma}$$

– Useful for convergence

- Prediction of immediate and future rewards differentially recruits cortico-basal ganglia loops



S.C. Tanaka, K. Doya, G. Okada, K. Ueda, Y. Okamoto, and S. Yamawaki. (2004). [Prediction of immediate and future rewards differentially recruits cortico-basal ganglia loops](#). Nature Neuroscience, 7(8): 887-893.

State-Value Function

- A state value function evaluates the policy in each state
- For a given stationary policy π , consider a state-action-reward sequence


$$s_t, a_t, r_{t+1}, s_{t+1}, a_{t+1}, r_{t+2}, s_{t+2} \dots$$

generated by

$$a_t \sim \pi(a \mid s_t), \quad s_{t+1}, r_{t+1} \sim p(s', r \mid s_t, a_t)$$

- The state value function is the expected return given by

$$V^\pi(s) \triangleq \mathbb{E}_\pi[G_t \mid s_t = s, a_{t:\infty} \sim \pi]$$



State-Value Function

- The value of a state s under a policy π is defined by

$$V^\pi(s) \triangleq \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s \right]$$

- $\mathbb{E}_{\pi}\{ \quad \}$ denotes the expectation which is taken over the probability distribution

$$\begin{aligned} &P(a_t, r_{t+1}, s_{t+1}, a_{t+1}, r_{t+2}, s_{t+2}, \dots \mid s_t = s) \\ &= \prod_{k=0}^{\infty} \pi(a_{t+k} \mid s) P_T(s_{t+k+1}, r_{t+k+1} \mid s_{t+k}, a_{t+k}) \end{aligned}$$


- Therefore,

$$V^\pi(s) = \sum_{a_t, r_{t+1}, s_{t+1}, \dots} \Pr(a_t, r_{t+1}, s_{t+1}, \dots \mid s_t = s) \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \right]$$

Example of State-Value Function

- Task: to get the battery pack while avoiding collisions with an obstacle
 - Positive reward for successful battery catching
 - Negative reward for collisions with obstacles

The goal state has the highest value




state value function

Action-Value Function

- An action-value function used in the bandit problem is extended to a function of state and action
- It is the expected return starting from s , taking the action a , and **thereafter following policy π**
 - $s_{t+1}, r_{t+1} \sim p(s', r \mid s_t, a_t), a_{t+1} \sim \pi(a \mid s_{t+1})$
- The state-action value function is the expected return given by

$$Q^\pi(s, a) \triangleq \mathbb{E}_\pi[G_t \mid s_t = s, a_t = a]$$



Action-Value Function

- The value of a state s under a policy π is defined by

$$Q^\pi(s, a) \triangleq \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s, a_t = a \right]$$

- $\mathbb{E}_\pi\{ \quad \}$ denotes the expectation which is taken over the probability distribution

$$\begin{aligned} &P(r_{t+1}, s_{t+1}, a_{t+1}, s_{t+2}, \dots \mid s_t = s, a_t = a) \\ &= \prod_{k=0}^{\infty} p(s_{t+k+1}, r_{t+k+1} \mid s_{t+k}, a_{t+k}) \pi(a_{t+k+1} \mid s_{t+k+1}) \end{aligned}$$

- Therefore,


$$Q^\pi(s, a) = \sum_{r_{t+1}, s_{t+1}, a_{t+1}, \dots} P(r_{t+1}, s_{t+1}, \dots \mid s_t = s, a_t = a) \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \right]$$

Relation between V^π and Q^π

- Relation between $V^\pi(s)$ and $Q^\pi(s, a)$

$$\begin{aligned} V^\pi(s) &= \sum_{a \in \mathcal{A}} \pi(a | s) Q^\pi(s, a) \\ &= \mathbb{E}_\pi[Q^\pi(s, a)] \end{aligned}$$

- We have already used this relation




Advantage function

- Advantage function

$$A_{\pi}(s, a) \triangleq q_{\pi}(s, a) - V^{\pi}(s)$$

- From the definition,

$$\begin{aligned}\mathbb{E}_{\pi}[A_{\pi}(s, a)] &= \mathbb{E}_{\pi}[Q^{\pi}(s, a)] - V^{\pi}(s) \\ &= 0\end{aligned}$$



Bellman (Expectation) Equation

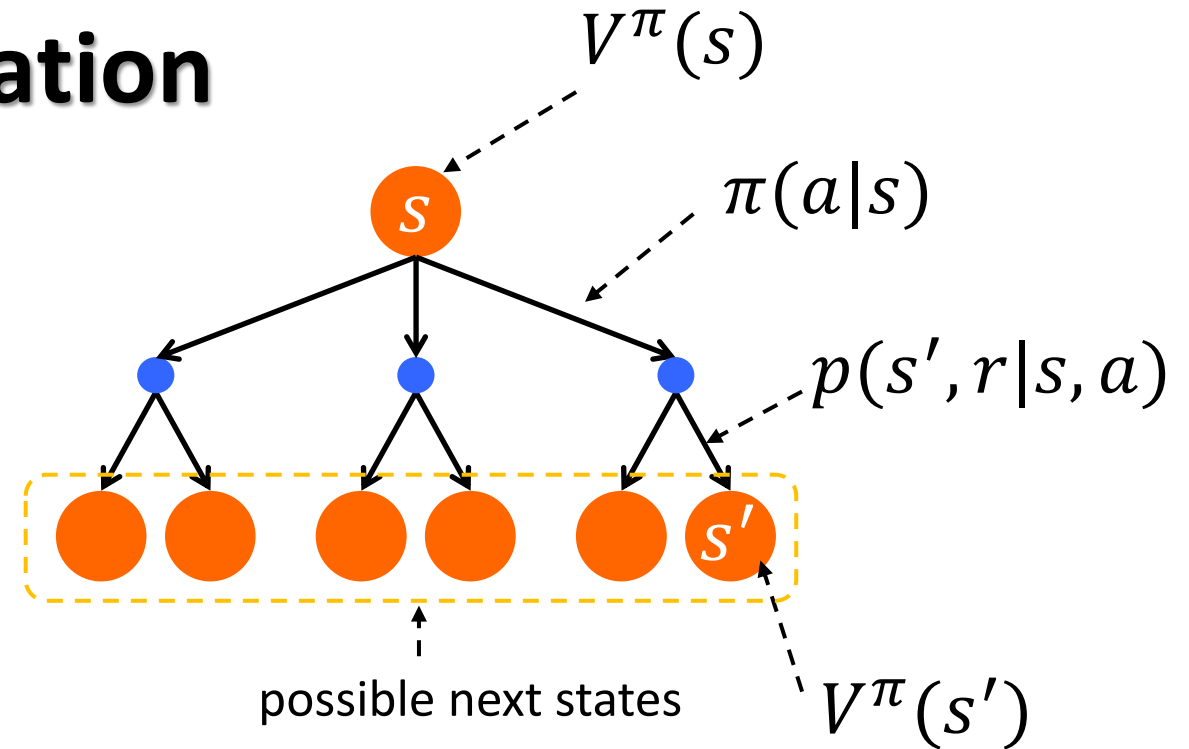
- Recursive relationship
- **Linear** with respect to V^π
- For any policy π and any state s ,

$$V^\pi(s) = \mathbb{E}_\pi[G_t \mid s_t = s]$$

$$= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid s_t = s]$$

$$= \sum_a \pi(a \mid s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma \mathbb{E}_\pi[G_{t+1} \mid s_{t+1} = s']]$$

$$= \sum_a \pi(a \mid s) \sum_{s', r} P(s', r \mid s, a) [r + \gamma V^\pi(s')]$$




Bellman (Expectation) Equation

- Recursive relationships for Q^π
- **Linear** with respect to Q^π
- For any policy π and any state-action pair (s, a) ,

$$Q^\pi(s, a) = \mathbb{E}_\pi[G_t \mid s_t = s, a_t = a]$$

$$= \sum_{s', r} p(s', r \mid s, a) [r + \gamma \mathbb{E}_\pi[G_{t+1} \mid s_{t+1} = s', a_{t+1} = a']]$$

$$= \sum_{s', r} P(s', r \mid s, a) \left[r + \gamma \sum_{a'} \pi(a' \mid s') Q^\pi(s', a') \right]$$




Relations between V^π and Q^π

- Using the previous equations, we obtain

$$Q^\pi(s, a) = \sum_{s', r} P(s', r \mid s, a) \left[r + \gamma \sum_{a'} \pi(a' \mid s') Q^\pi(s', a') \right]$$

$$= \sum_{s', r} p(s', r \mid s, a) [r + \gamma V^\pi(s')]$$



Optimal Value Functions

- Reminder: The goal of RL is to find an optimal policy π_*
- V^* is an optimal state value function for π^* defined by

$$V^*(s) \triangleq \max_{\pi} V^{\pi}(s) \quad \forall s \in \mathcal{S}$$

- Similarly, an optimal state-action value function is defined by

$$Q^*(s, a) \triangleq \max_{\pi} Q^{\pi}(s, a) \quad \forall s \in \mathcal{S}, \forall a \in \mathcal{A}$$

Optimal Value Functions

Results

Here, p_0 is the probability to select the action a_0 at the state s_0 . Then, the probability to select a_1 at s_0 is given by $1 - p_0$. Similarly, p_1 is the probability to select a_0 at s_1 . The first array shows the state-value function, $V^\pi(s)$. The second 2-D array shows the action-value function, $Q^\pi(s, a)$.

```
env = TwoStageMDPEnv()
interact(compute_value, env=fixed(env), p0=(0, 1, 0.1), p1=(0, 1, 0.1), p2=(0, 1, 0.1), p3=(0, 1, 0.1), p4=(0, 1, 0.1))
```



p0 0.50
p1 0.50
p2 0.50
p3 0.50
p4 0.50

} $\pi(a_0 | s)$

(array([[7.6125],
[5.],
[5.25],
[12.25],
[5.25],
[0.]]), array([[5.075, 10.15],
[7. , 3.],
[3. , 7.5],
[17.5 , 7.],
[3. , 7.5],
[0. , 0.]]))

} $V^\pi(s)$

} $Q^\pi(s, a)$

<function __main__.compute_value(env, p0=0.5, p1=0.5, p2=0.5, p3=0.5, p4=0.5)>

Try finding the optimal policy

[4]

Theorem

- A state-value function defines a partial ordering over policies

$$\pi \geq \pi' \quad \text{if} \quad V^\pi(s) \geq V^{\pi'}(s), \quad \text{for all } s$$

Theorem

- For any Markov Decision Process

1. There exists an optimal policy π^* that is better than or equal to all other policies, $\pi^* \geq \pi$ for all π
2. All optimal policies achieve the optimal value function,

$$V^{\pi^*}(s) = V^*(s)$$

3. All optimal policies achieve the optimal state-action value function

$$\underbrace{Q^{\pi^*}(s, a)}_{\text{action-value function of the optimal policy}} = \underbrace{Q^*(s, a)}_{\text{optimal action-value function}}$$

action-value function of the
optimal policy

optimal action-value function

Bellman Optimality Equation for V^*

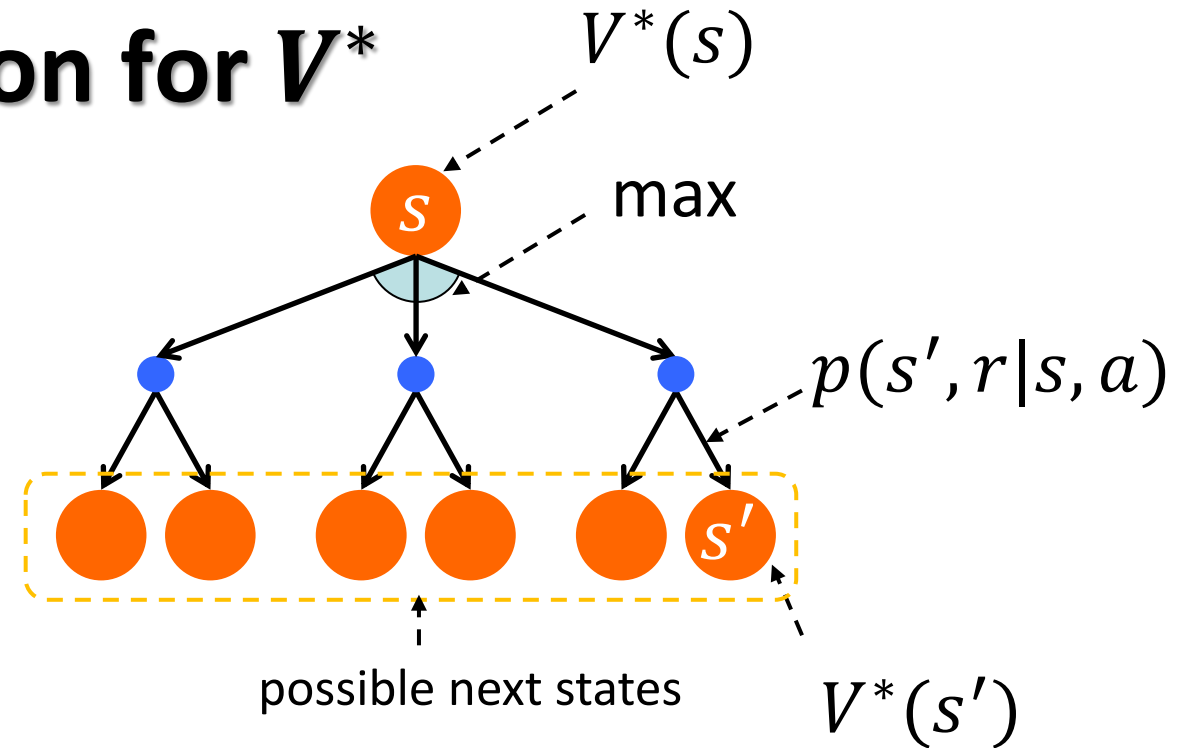
- Recursive relationships
- **Nonlinear** with respect to V^*
- For any state s ,

$$V^*(s) = \max_a Q^*(s, a)$$

$$= \max_a \mathbb{E}_{\pi^*}[G_t \mid s_t = s, a_t = a] = \max_a \mathbb{E}_{\pi^*}[r_{t+1} + \gamma G_t \mid s_t = s, a_t = a]$$

$$= \max_a \mathbb{E}[r_{t+1} + \gamma V^*(s_{t+1}) \mid s_t = s, a_t = a]$$

$$= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma V^*(s')]$$



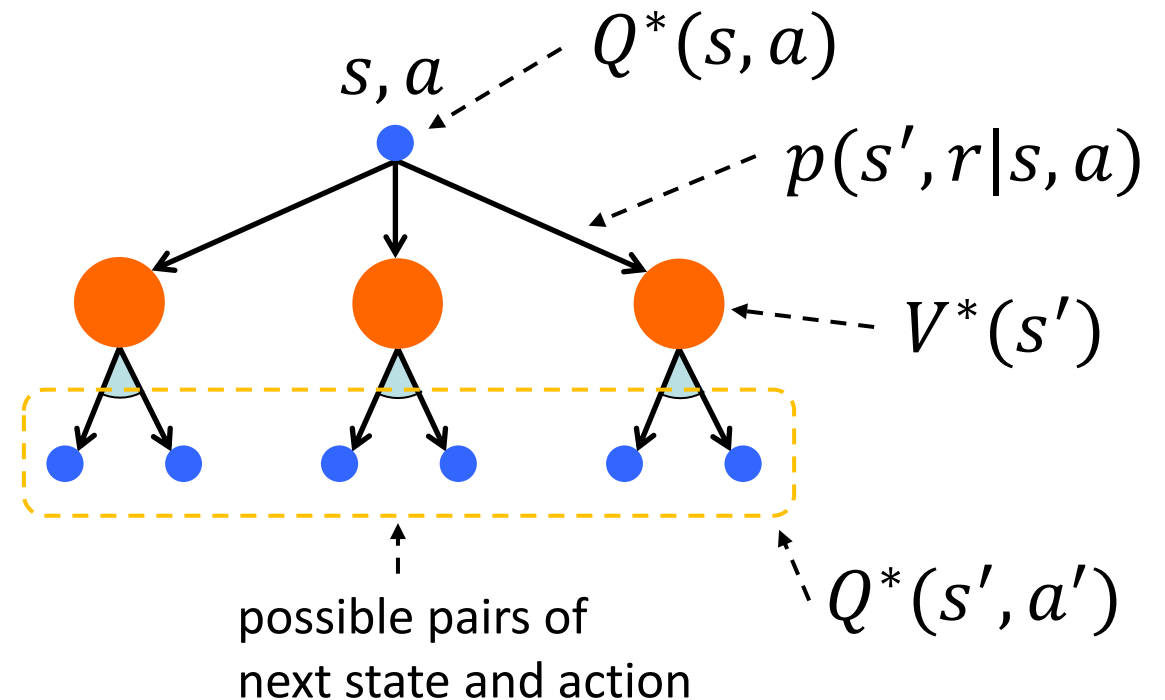
Bellman Optimality Equation for Q^*

- Similarly, the Bellman optimality equation is given by

$$Q^*(s, a) = \mathbb{E}[r_{t+1} + \gamma \max_{a'} Q^*(s_{t+1}, a') \mid s_t = s, a_t = a]$$

$$= \sum_{s', r} p(s', r \mid s, a) [r + \gamma \max_{a'} Q^*(s', a')]$$

- This is also a nonlinear equation with respect to Q^* due to the max operator



Optimal Value Functions

- There exists the following relationships:

$$V^*(s) = \mathbb{E} \left[r_{t+1} + \gamma \max_a Q^*(s_{t+1}, a) \mid s_t = s \right]$$

$$Q^*(s, a) = \mathbb{E}[r_{t+1} + \gamma V^*(s_{t+1}) \mid s_t = s, a_t = a]$$

Summary

- Bellman expectation equation

$$V^{\pi}(s) = \sum_a \pi(a | s) \sum_{s', r} p(s', r | s, a) [r + \gamma V^{\pi}(s')]$$

$$Q^{\pi}(s, a) = \sum_{s', r} p(s', r | s, a) \left[r + \gamma \sum_{a'} \pi(a' | s') Q^{\pi}(s', a') \right]$$

- Bellman optimality equation

$$V^*(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V^*(s')]$$

$$Q^*(s, a) = \sum_{s', r} p(s', r | s, a) [r + \gamma \max_{a'} Q^*(s', a')]$$